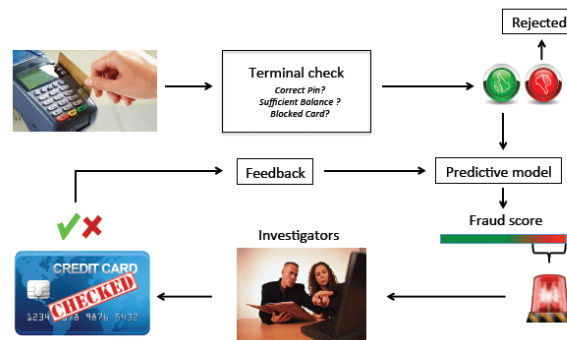


Credit Card Fraud Detection



Problem Statement:

- For many banks, retaining high profitable customers is the number one business goal. Banking fraud, however, poses a significant threat to this goal for different banks. In terms of substantial financial losses, trust and credibility, this is a concerning issue to both banks and customers alike.
- In the banking industry, credit card fraud detection using machine learning and deep learning is not only a trend but a necessity for them to put proactive monitoring and fraud prevention mechanisms in place. Machine learning is helping these institutions to reduce time-consuming manual reviews, costly chargebacks and fees as well as denials of legitimate transactions.
- In this project we will detect fraudulent credit card transactions with the help of Machine learning and deep learning models.
- We will analyse customer-level data that has been collected and analysed during a research collaboration of Worldline and the Machine Learning Group.

Theory about Credit Card Fraud

- What is Credit Card Fraud Detection?**
 - Credit card fraud is a term that has been coined for unauthorized access of payment cards like credit cards or debit cards to pay for using services or goods
 - Hackers or fraudsters may obtain the confidential details of the card from unsecured websites. When a fraudster compromises an individual's credit/debit card, everyone involved in the process suffers, right from the individual whose confidential data has been leaked to the businesses (generally banks) who issue the credit card and the merchant who is finalizing the transaction with purchase
 - This makes it extremely essential to identify the fraudulent transactions at the onset. Financial institutions and businesses like e-commerce are taking firm steps to flag the fraudsters entering the system. Various advanced machine learning technologies are at play, assessing every transaction and stemming the fraud users in its nip using behavioral data and transaction patterns
 - The process of automatically differentiating between fraudulent and genuine users is known as "credit card fraud detection"
- How does Credit Card Fraud work?**
 - A credit card is one of the most used financial products to make online purchases and payments such as gas, groceries, TVs, traveling, shopping bills, and so on because of the non-availability of funds at that instance. Credit cards are of most value that provide various benefits in the form of points while using them for different transactions. There are several categories of credit card fraud that are prevalent in today's time:
 - Lost/Stolen cards:** People steal credit cards from the mail and use them illegally on behalf of the owner. The process of blocking credit cards that have been stolen and re-issuing them is a hassle for both customers and credit card companies. Some financial institutions keep the credit cards blocked until it is verified that the rightful owner has received the card.
 - Card Abuse:** The customer buys goods and items on the credit card but has no intention to pay back the amount charged by the bank for the same. These customers stop answering the calls as the deadline to settle the dues approaches. Sometimes they even declare bankruptcy—this type of fraud results in losses of millions every year.
 - Identity Theft:** The customers apply illegitimate information, and they might even steal the details of a genuine customer to apply for a credit card and then misuse it. In such cases, even card blocking can not stop the credit card from falling into the wrong hands.
 - Merchant Abuse:** Some merchants show illegal transactions (that never occurred) for money laundering. For performing these illicit transactions, legal information of genuine credit card users is stolen to generate replicas of the cards and use it for illegal work.
 - Many traditional old-school techniques have been used since time in-memorial for credit card fraud detection like CVV verification, geolocation tracking, IP Address verification, etc. But over time, the criminals are using more advanced techniques to commit crimes, and it is impossible to prevent them all using only traditional methods. Millions of transactions are processed every second in today's world, which takes it beyond human intelligence to process all the data to identify the behavioral patterns of the fraudsters. This is where credit card fraud detection using machine learning plays a vital role.
 - Financial institutions increasingly depend upon automated machine learning systems to make intelligent decisions and protect businesses against substantial losses. These measures play a significant role in reducing the risk while doing online transactions. Like humans, machine learning algorithms learn from past transaction data and use that information to analyze future transactions with the same lens. While machines might not be as intelligent as humans are and might need some supervision on top of it, the advantage lies in the speed of data processing and computation. Also, machines can identify and remember more patterns in vast volumes of data compared to humans. Generally, these algorithms are known as anomaly detection. Let us delve into details in the next section.

Credit Card Fraud Detection using Machine Learning can be done using

1. Unsupervised Learning -

- Machine Learning Algorithms such as Isolation Forest, One-class SVM, LOF, etc., do not require labeled data for training the model. They identify patterns in the data and try to group the data points based on observed similarities in patterns.

2. Supervised Learning

- Machine Learning and Deep Learning Algorithms such as Ensemble Models (RandomForest, XGBoost, LightGBM, etc.), KNN, Neural Networks, Autoencoders, etc. These algorithms are trained on labeled data, and the model learns to predict the labels for the unseen data. Labeled data can be expensive to gather.

Challenges in Credit Card Fraud Detection

- The challenges involved in credit card fraud detection project is primarily the data itself. The data is heavily imbalanced, i.e., the count of data labeled as fraudulent is way less than the data labeled as non-fraudulent data. This makes it extremely tricky to train the model as it tends to overfit for the majority class and underfit for the minority class. Techniques like oversampling, undersampling, cost-sensitive learning, etc. can be used to deal with this. The metrics used for the final model are different from standard evaluation metrics of accuracy, AUC-ROC, etc.
- Another prevalent faced challenge is the quality and quantity of the data. The startups in the early stage do not have much user history data to train extensive models, which makes it difficult to train a robust fraud detection model. A temporary solution to this problem can be sourcing data from an external third party, like scores from credit bureaus.

DATASET DESCRIPTION

- The dataset contains 284,807 transactions among which there are 492 i.e., 0.172% transactions are fraudulent transactions
- It also contains transactions made by a cardholder in 2 days in month of september 2013
- This dataset is highly unbalanced. Due to security reasons, most of the features in the dataset are transformed using principal component analysis (PCA). V1, V2, V3,...., V28 are PCA applied features and rest features include 'time', 'amount' and 'class' are non-PCA applied features

Table of Contents

- 1. Importing dependencies
- 2. Exploratory data analysis
- 3. Splitting the data into train & test data
- 4. Model Building
 - Perform cross validation with RepeatedKfold
 - Perform cross validation with StratifiedKfold
 - RandomOverSampler with StratifiedKfold Cross Validation
 - Oversampling with SMOTE Oversampling
 - Oversampling with ADASYN Oversampling
- 5. Hyperparameter Tuning
- 6. Conclusion

Importing Dependencies

```
In [340]: # Importing the libraries
import numpy as np
import pandas as pd
import time

import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

from scipy import stats
from scipy.stats import norm, skew
from scipy.special import boxcox1p
from scipy.stats import boxcox_normmax

from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler

import sklearn
from sklearn import metrics
from sklearn.metrics import roc_curve, auc, roc_auc_score
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import average_precision_score, precision_recall_curve

from sklearn.model_selection import train_test_split
from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV

from sklearn.linear_model import Ridge, Lasso, LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegressionCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import RandomForestClassifier
import xgboost as xgb
from xgboost import XGBClassifier
from xgboost import plot_importance
from sklearn.ensemble import AdaBoostClassifier

from tensorflow.keras import Sequential
import tensorflow as tf
from tensorflow.keras.layers import Dense, Dropout

# To ignore warnings
import warnings
warnings.filterwarnings("ignore")
```

Exploratory data analysis

```
In [341]: # Mounting the google drive
from google.colab import drive
drive.mount('/content/drive')
import os
os.getcwd()
path = "/content/drive/MyDrive/Projects/Credit Card Defaulter"
os.chdir(path)
os.getcwd()

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

Out[341]: '/content/drive/MyDrive/Projects/Credit Card Defaulter'
```

```
In [342]: # Loading the data
df = pd.read_csv('creditcard.csv')
df = df.sample(n=50000)
df.head()
```

```
Out[342]:
```

| | Time | V1 | V2 | V3 | V4 | V5 | V6 | V7 | V8 | V9 | ... | V21 | V22 | V23 | V24 | V25 | V26 | V27 | V28 | Amount | Class |
|--------|----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----|-----------|----------|-----------|----------|-----------|-----------|-----------|-----------|--------|-------|
| 192316 | 129633.0 | -0.541013 | 1.107931 | -1.774693 | -0.708025 | -0.456628 | -0.295740 | 2.360015 | -0.293290 | -0.597604 | ... | 0.100783 | 0.308628 | -0.177781 | 0.642998 | 0.037229 | 0.494459 | -0.390304 | -0.202842 | 300.00 | 0 |
| 91715 | 63596.0 | -0.627415 | 1.123128 | 1.560804 | -0.076211 | 0.223046 | -0.904104 | 0.938718 | -0.155628 | -0.392705 | ... | 0.067629 | 0.094912 | -0.176009 | 0.359548 | -0.058737 | -0.666741 | -0.019712 | 0.135237 | 4.65 | 0 |
| 140052 | 83511.0 | -0.916656 | 0.996912 | 2.114384 | 1.082337 | -1.104885 | -0.049506 | -0.411277 | 0.755500 | 0.369418 | ... | -0.006409 | 0.077001 | -0.108486 | 0.345586 | -0.031918 | -0.303443 | 0.275835 | 0.123201 | 9.99 | 0 |
| 13800 | 24466.0 | -0.886470 | -0.126264 | 3.551005 | 3.751230 | -0.563369 | 1.141458 | -0.889331 | 0.234025 | 2.022399 | ... | -0.058542 | 0.915166 | 0.376580 | 0.375107 | -0.396255 | 0.368435 | 0.028081 | -0.129019 | 3.80 | 0 |
| 55873 | 47165.0 | 1.223111 | -0.942174 | 0.196413 | -0.420355 | -1.078497 | -0.451728 | -0.572279 | 0.023307 | -0.216677 | ... | 0.099471 | 0.171263 | -0.112221 | 0.065022 | 0.518764 | -0.116556 | -0.004219 | 0.012913 | 80.44 | 0 |

5 rows × 31 columns

```
In [343]: # Checking the shape
df.shape
```

```
Out[343]: (50000, 31)
```

```
In [344]: # Checking the datatypes and null/non-null distribution
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 50000 entries, 192316 to 240041
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    Time      50000 non-null   float64
1    V1         50000 non-null   float64
2    V2         50000 non-null   float64
3    V3         50000 non-null   float64
4    V4         50000 non-null   float64
5    V5         50000 non-null   float64
6    V6         50000 non-null   float64
7    V7         50000 non-null   float64
8    V8         50000 non-null   float64
9    V9         50000 non-null   float64
10   V10        50000 non-null   float64
11   V11        50000 non-null   float64
12   V12        50000 non-null   float64
13   V13        50000 non-null   float64
14   V14        50000 non-null   float64
15   V15        50000 non-null   float64
16   V16        50000 non-null   float64
17   V17        50000 non-null   float64
18   V18        50000 non-null   float64
19   V19        50000 non-null   float64
20   V20        50000 non-null   float64
21   V21        50000 non-null   float64
22   V22        50000 non-null   float64
23   V23        50000 non-null   float64
24   V24        50000 non-null   float64
25   V25        50000 non-null   float64
26   V26        50000 non-null   float64
27   V27        50000 non-null   float64
28   V28        50000 non-null   float64
29   Amount    50000 non-null   float64
30   Class     50000 non-null   int64
dtypes: float64(30), int64(1)
memory usage: 12.2 MB
```

```
In [345]: # Checking distribution of numerical values in the dataset
df.describe()
```

Out[345]:

| | Time | V1 | V2 | V3 | V4 | V5 | V6 | V7 | V8 | V9 | ... | V21 | V22 | V23 | V24 | |
|-------|---------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|-----|--------------|--------------|--------------|--------------|--------------|
| count | 50000.000000 | 50000.000000 | 50000.000000 | 50000.000000 | 50000.000000 | 50000.000000 | 50000.000000 | 50000.000000 | 50000.000000 | 50000.000000 | ... | 50000.000000 | 50000.000000 | 50000.000000 | 50000.000000 | 50000.000000 |
| mean | 95027.093240 | 0.006889 | 0.000014 | 0.002052 | 0.006169 | 0.003853 | 0.000950 | 0.004717 | 0.004668 | 0.008224 | ... | 0.003217 | -0.002990 | 0.000050 | -0.002233 | -0.000000 |
| std | 47608.097497 | 1.940516 | 1.625471 | 1.504460 | 1.410004 | 1.364510 | 1.322307 | 1.212862 | 1.170605 | 1.091409 | ... | 0.726534 | 0.725087 | 0.601028 | 0.606238 | 0.516000 |
| min | 0.000000 | -40.470142 | -42.172688 | -31.103685 | -5.683171 | -40.427726 | -23.496714 | -31.197329 | -50.943369 | -9.462573 | ... | -22.665685 | -8.887017 | -23.222016 | -2.740677 | -4.930000 |
| 25% | 54267.750000 | -0.912608 | -0.597673 | -0.883402 | -0.841195 | -0.683034 | -0.763299 | -0.549844 | -0.204852 | -0.636216 | ... | -0.229721 | -0.546591 | -0.160259 | -0.356284 | -0.316000 |
| 50% | 85203.000000 | 0.017080 | 0.072429 | 0.176216 | -0.011505 | -0.053596 | -0.270196 | 0.042177 | 0.022762 | -0.046710 | ... | -0.029802 | 0.004218 | -0.009554 | 0.041472 | 0.016000 |
| 75% | 139631.500000 | 1.315053 | 0.808100 | 1.024045 | 0.759097 | 0.609536 | 0.399615 | 0.573382 | 0.327584 | 0.607525 | ... | 0.187433 | 0.531716 | 0.149508 | 0.437811 | 0.340000 |
| max | 172792.000000 | 2.451888 | 16.713389 | 3.920275 | 12.699542 | 34.099309 | 23.917837 | 44.054461 | 20.007208 | 10.392889 | ... | 27.202839 | 7.220158 | 15.426351 | 3.998294 | 4.820000 |

8 rows × 31 columns


```
In [346]: # Checking the class distribution of the target variable
df['Class'].value_counts()
```

Out[346]:

```
0    49916
1      84
Name: Class, dtype: int64
```

```
In [347]: # Checking the class distribution of the target variable in percentage
df['Class'].value_counts(normalize = True).plot.pie()
```

Out[347]: <matplotlib.axes._subplots.AxesSubplot at 0x7f53a19b95e0>




```
In [348]: # Checking the correlation
corr = df.corr()
corr
```

Out[348]:

| | Time | V1 | V2 | V3 | V4 | V5 | V6 | V7 | V8 | V9 | ... | V21 | V22 | V23 | V24 | V25 | V26 | V27 | V28 | Amount | Cl |
|--------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|--------|
| Time | 1.000000 | 0.118292 | -0.012198 | -0.421144 | -0.112760 | 0.177543 | -0.058231 | 0.082175 | -0.032188 | -0.005067 | ... | 0.046404 | 0.147353 | 0.051976 | -0.017890 | -0.239786 | -0.044932 | -0.006478 | -0.011512 | -0.012154 | -0.016 |
| V1 | 0.118292 | 1.000000 | -0.004389 | -0.009109 | 0.011063 | -0.007697 | -0.001057 | 0.004101 | -0.011941 | 0.005733 | ... | -0.014567 | 0.007841 | 0.040755 | -0.004787 | 0.010659 | -0.006836 | -0.052209 | 0.047343 | -0.221128 | -0.107 |
| V2 | -0.012198 | -0.004389 | 1.000000 | 0.001119 | 0.009434 | 0.007904 | 0.003128 | 0.022083 | 0.005633 | -0.006709 | ... | 0.022919 | -0.002429 | 0.059792 | -0.000121 | 0.014421 | 0.000852 | -0.037145 | 0.059617 | -0.526473 | 0.085 |
| V3 | -0.421144 | -0.009109 | 0.001119 | 1.000000 | 0.014623 | -0.009876 | 0.003400 | -0.003018 | -0.008532 | 0.000546 | ... | -0.016443 | 0.015221 | 0.014239 | -0.000074 | -0.002764 | -0.005301 | -0.011037 | 0.013405 | -0.206868 | -0.187 |
| V4 | -0.112760 | 0.011063 | 0.009434 | 0.014623 | 1.000000 | 0.003917 | -0.006781 | 0.004598 | -0.003612 | 0.003839 | ... | -0.002481 | -0.001170 | -0.016054 | 0.000745 | -0.001209 | -0.001193 | 0.009019 | -0.010950 | 0.085269 | 0.118 |
| V5 | 0.177543 | -0.007697 | 0.007904 | -0.009876 | 0.003917 | 1.000000 | 0.006577 | 0.004568 | -0.007636 | -0.004547 | ... | -0.018828 | 0.013145 | -0.002647 | 0.000292 | -0.011059 | 0.000032 | 0.013777 | -0.017210 | -0.385417 | -0.098 |
| V6 | -0.058231 | -0.001057 | 0.003128 | 0.003400 | -0.006781 | 0.006577 | 1.000000 | -0.002341 | 0.019124 | -0.001879 | ... | 0.013933 | 0.000243 | 0.002962 | -0.009796 | 0.001131 | -0.003912 | -0.015692 | 0.013130 | 0.210204 | -0.040 |
| V7 | 0.082175 | 0.004101 | 0.022083 | -0.003018 | 0.004598 | 0.004568 | -0.002341 | 1.000000 | -0.018584 | -0.020129 | ... | -0.020720 | 0.009716 | -0.015913 | 0.003073 | -0.002654 | -0.003813 | -0.026919 | 0.009925 | 0.390978 | -0.173 |
| V8 | -0.032188 | -0.011941 | 0.005633 | -0.008532 | -0.003612 | -0.007636 | 0.019124 | -0.018584 | 1.000000 | -0.005950 | ... | -0.054987 | 0.033091 | 0.020526 | -0.008456 | 0.004981 | -0.006505 | -0.020052 | 0.008887 | -0.106148 | 0.031 |
| V9 | -0.005067 | 0.005733 | -0.006709 | 0.000546 | 0.003839 | -0.004547 | -0.001879 | -0.020129 | -0.005950 | 1.000000 | ... | -0.002835 | 0.009277 | -0.006674 | -0.002013 | -0.011032 | 0.000511 | 0.013312 | -0.021642 | -0.043101 | -0.083 |
| V10 | 0.030502 | 0.004307 | -0.000151 | 0.000223 | 0.015494 | -0.005913 | 0.003096 | -0.024081 | -0.017924 | -0.030983 | ... | -0.017353 | -0.000331 | 0.002377 | 0.003696 | -0.003346 | 0.004854 | 0.012708 | -0.019358 | -0.099478 | -0.197 |
| V11 | -0.250578 | 0.000665 | -0.000712 | -0.001009 | -0.002416 | 0.001904 | -0.003124 | -0.000269 | 0.004796 | -0.000142 | ... | 0.002300 | -0.006485 | -0.000117 | -0.002833 | 0.002569 | 0.000449 | 0.002089 | 0.002802 | -0.002800 | 0.144 |
| V12 | 0.125820 | 0.000583 | -0.002155 | -0.011736 | 0.002873 | -0.000296 | 0.002972 | -0.010803 | -0.006058 | -0.012543 | ... | 0.002496 | 0.003197 | 0.001026 | -0.004147 | 0.000852 | -0.007224 | 0.012033 | -0.004487 | -0.008161 | -0.234 |
| V13 | -0.062962 | 0.001689 | -0.003999 | 0.002939 | -0.001558 | -0.001034 | 0.003397 | -0.005260 | 0.000664 | 0.004843 | ... | -0.008727 | 0.001335 | -0.004190 | -0.001789 | -0.004900 | 0.002078 | -0.000077 | -0.004109 | 0.006545 | 0.000 |
| V14 | -0.101219 | -0.003645 | 0.000807 | 0.000297 | 0.003207 | 0.002666 | 0.004044 | 0.002015 | -0.004924 | 0.002770 | ... | -0.008612 | 0.001467 | 0.003836 | 0.000799 | 0.000844 | 0.001969 | -0.005305 | 0.010606 | 0.032959 | -0.270 |
| V15 | -0.182811 | 0.001768 | 0.002957 | -0.004989 | 0.000194 | -0.003227 | 0.001456 | -0.000049 | -0.003204 | -0.009113 | ... | -0.000029 | -0.001682 | -0.006076 | 0.002694 | 0.008768 | 0.000749 | 0.011544 | -0.006817 | -0.000919 | -0.005 |
| V16 | 0.015188 | 0.002290 | 0.002878 | -0.001924 | 0.003615 | 0.002571 | -0.008992 | -0.009611 | 0.001285 | 0.000524 | ... | -0.006569 | -0.001318 | -0.003929 | 0.006953 | -0.000389 | 0.010345 | 0.003444 | -0.010813 | -0.011884 | -0.179 |
| V17 | -0.075270 | 0.003702 | -0.001134 | -0.003380 | 0.010248 | -0.000359 | -0.001768 | 0.001545 | -0.011023 | -0.002319 | ... | -0.021812 | 0.007780 | 0.004943 | -0.001321 | -0.000501 | 0.003339 | -0.009543 | 0.005462 | 0.010560 | -0.304 |
| V18 | 0.095193 | 0.006327 | 0.003657 | -0.006155 | -0.007904 | 0.002368 | -0.006642 | -0.003419 | -0.003290 | 0.001512 | ... | -0.011524 | 0.009922 | -0.006112 | -0.002160 | -0.004865 | 0.003844 | -0.002958 | -0.001955 | 0.030029 | -0.105 |
| V19 | 0.028739 | -0.002696 | -0.003370 | -0.002532 | 0.006034 | 0.002800 | -0.011828 | -0.004611 | -0.010370 | 0.003776 | ... | 0.000664 | -0.009657 | -0.007034 | 0.003117 | -0.003531 | -0.007377 | 0.009047 | -0.004321 | -0.051108 | 0.034 |
| V20 | -0.051428 | 0.019607 | 0.021118 | 0.004534 | -0.006939 | -0.025632 | 0.013653 | 0.037741 | 0.014759 | -0.000393 | ... | -0.004366 | 0.012152 | 0.029424 | 0.005008 | 0.010408 | 0.005749 | -0.036387 | 0.044919 | 0.332523 | 0.014 |
| V21 | 0.046404 | -0.014567 | 0.022919 | -0.016443 | -0.002481 | -0.018828 | 0.013933 | -0.020720 | -0.054987 | -0.002835 | ... | 1.000000 | 0.020952 | 0.017877 | -0.001753 | 0.004629 | -0.004908 | -0.017734 | 0.022524 | 0.104362 | 0.066 |
| V22 | 0.147353 | 0.007841 | -0.002429 | 0.015221 | -0.001170 | 0.013145 | 0.000243 | 0.009716 | 0.033091 | 0.009277 | ... | 0.020952 | 1.000000 | -0.004213 | 0.001411 | -0.001968 | -0.001008 | -0.002465 | 0.003759 | -0.068103 | -0.012 |
| V23 | 0.051976 | 0.040755 | 0.059792 | 0.014239 | -0.016054 | -0.002647 | 0.002962 | -0.015913 | 0.020526 | -0.006674 | ... | 0.017877 | -0.004213 | 1.000000 | -0.003951 | -0.021767 | -0.003337 | -0.055086 | 0.046792 | -0.159856 | 0.003 |
| V24 | -0.017890 | -0.004787 | -0.000121 | -0.000074 | 0.000745 | 0.000292 | -0.009796 | 0.003073 | -0.008456 | -0.002013 | ... | -0.001753 | 0.001411 | -0.003951 | 1.000000 | 0.005079 | 0.004826 | 0.008047 | -0.001192 | 0.007971 | -0.008 |
| V25 | -0.239786 | 0.010659 | 0.014421 | -0.002764 | -0.001209 | -0.011059 | 0.001131 | -0.002654 | 0.004981 | -0.011032 | ... | 0.004629 | -0.001968 | -0.021767 | 0.005079 | 1.000000 | 0.003479 | -0.013041 | -0.010587 | -0.047447 | 0.009 |
| V26 | -0.044932 | -0.006836 | 0.000852 | -0.005301 | -0.001193 | 0.000032 | -0.003912 | -0.003813 | -0.006505 | 0.000511 | ... | -0.004908 | -0.001008 | -0.003337 | 0.004826 | 0.003479 | 1.000000 | -0.007550 | 0.004986 | 0.001026 | 0.002 |
| V27 | -0.006478 | -0.052209 | -0.037145 | -0.011037 | 0.009019 | 0.013777 | -0.015692 | -0.026919 | -0.020052 | 0.013312 | ... | -0.017734 | -0.002465 | -0.055086 | 0.008047 | -0.013041 | -0.007550 | 1.000000 | 0.073469 | 0.038908 | 0.026 |
| V28 | -0.011512 | 0.047343 | 0.059617 | 0.013405 | -0.010950 | -0.017210 | 0.013130 | 0.009925 | 0.008887 | -0.021642 | ... | 0.022524 | 0.003759 | 0.046792 | -0.001192 | -0.010587 | 0.004986 | 0.073469 | 1.000000 | 0.010640 | 0.005 |
| Amount | -0.012154 | -0.221128 | -0.526473 | -0.206868 | 0.085269 | -0.385417 | 0.210204 | 0.390978 | -0.106148 | -0.043101 | ... | 0.104362 | -0.068103 | -0.159856 | 0.007971 | -0.047447 | 0.001026 | 0.038908 | 0.010640 | 1.000000 | 0.008 |
| Class | -0.016079 | -0.107550 | 0.085935 | -0.187882 | 0.118884 | -0.098504 | -0.040388 | -0.173761 | 0.031241 | -0.083785 | ... | 0.066410 | -0.012237 | 0.003708 | -0.008032 | 0.009835 | 0.002806 | 0.026460 | 0.005176 | 0.008788 | 1.000 |

```
In [349]: # Checking the correlation in heatmap
plt.figure(figsize=(22,18))

sns.heatmap(corr, cmap="coolwarm", annot=True)
plt.show()
```




Here we will observe the distribution of our classes

```
In [350]: # Checking the % distribution of normal vs fraud
classes=df['Class'].value_counts()
normal_share=classes[0]/df['Class'].count()*100
fraud_share=classes[1]/df['Class'].count()*100

print(normal_share)
print(fraud_share)

99.832
0.168
```

```
In [351]: # Create a bar plot for the number and percentage of fraudulent vs non-fraudulent transactions
plt.figure(figsize=(15,5))
sns.countplot(df['Class'])
plt.title("Class Count", fontsize=18)
plt.xlabel("Record counts by class", fontsize=15)
plt.ylabel("Count", fontsize=15)
plt.show()
```



```
In [352]: # As time is given in relative fashion, we are using pandas.Timedelta which Represents a duration, the difference between two times or dates.
Delta_Time = pd.to_timedelta(df['Time'], unit='s')

#Create derived columns Mins and hours
df['Time_Day'] = (Delta_Time.dt.components.days).astype(int)
df['Time_Hour'] = (Delta_Time.dt.components.hours).astype(int)
df['Time_Min'] = (Delta_Time.dt.components.minutes).astype(int)
```

```
In [353]: # Drop unnecessary columns
# We will drop Time, as we have derived the Day/Hour/Minutes from the time column
df.drop('Time', axis = 1, inplace= True)
# We will keep only derived column hour, as day/minutes might not be very useful
df.drop(['Time_Day', 'Time_Min'], axis = 1, inplace= True)
```

Splitting the data into train & test data

```
In [354]: # Splitting the dataset into X and y
y= df['Class']
X = df.drop(['Class'], axis=1)
```

```
In [355]: # Checking some rows of X
X.head()
```

Out[355]:

| | V1 | V2 | V3 | V4 | V5 | V6 | V7 | V8 | V9 | V10 | ... | V21 | V22 | V23 | V24 | V25 | V26 | V27 | V28 | Amount | Time_Hou |
|--------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----|-----------|----------|-----------|----------|-----------|-----------|-----------|-----------|--------|----------|
| 192316 | -0.541013 | 1.107931 | -1.774693 | -0.708025 | -0.456628 | -0.295740 | 2.360015 | -0.293290 | -0.597604 | -0.513867 | ... | 0.100783 | 0.308628 | -0.177781 | 0.642998 | 0.037229 | 0.494459 | -0.390304 | -0.202842 | 300.00 | 1 |
| 91715 | -0.627415 | 1.123128 | 1.560804 | -0.076211 | 0.223046 | -0.904104 | 0.938718 | -0.155628 | -0.392705 | -0.595657 | ... | 0.067629 | 0.094912 | -0.176009 | 0.359548 | -0.058737 | -0.666741 | -0.019712 | 0.135237 | 4.65 | 1 |
| 140052 | -0.916656 | 0.996912 | 2.114384 | 1.082337 | -1.104885 | -0.049506 | -0.411277 | 0.755500 | 0.369418 | -0.500382 | ... | -0.006409 | 0.077001 | -0.108486 | 0.345586 | -0.031918 | -0.303443 | 0.275835 | 0.123201 | 9.99 | 2 |
| 13800 | -0.886470 | -0.126264 | 3.551005 | 3.751230 | -0.563369 | 1.141458 | -0.889331 | 0.234025 | 2.022399 | 0.262257 | ... | -0.058542 | 0.915166 | 0.376580 | 0.375107 | -0.396255 | 0.368435 | 0.028081 | -0.129019 | 3.80 | |
| 55873 | 1.223111 | -0.942174 | 0.196413 | -0.420355 | -1.078497 | -0.451728 | -0.572279 | 0.023307 | -0.216677 | 0.492314 | ... | 0.099471 | 0.171263 | -0.112221 | 0.065022 | 0.518764 | -0.116556 | -0.004219 | 0.012913 | 80.44 | 1 |

5 rows × 30 columns

```
In [356]: # Checking some rows of y
y.head()
```

```
Out[356]: 192316    0
          91715    0
          140052    0
          13800    0
          55873    0
          Name: Class, dtype: int64
```

```
In [357]: # Splitting the dataset using train test split
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=1, test_size=0.20)
```

Preserve X_test & y_test to evaluate on the test data once you build the model

```
In [358]: # Checking the spread of data post split
print(np.sum(y))
print(np.sum(y_train))
print(np.sum(y_test))

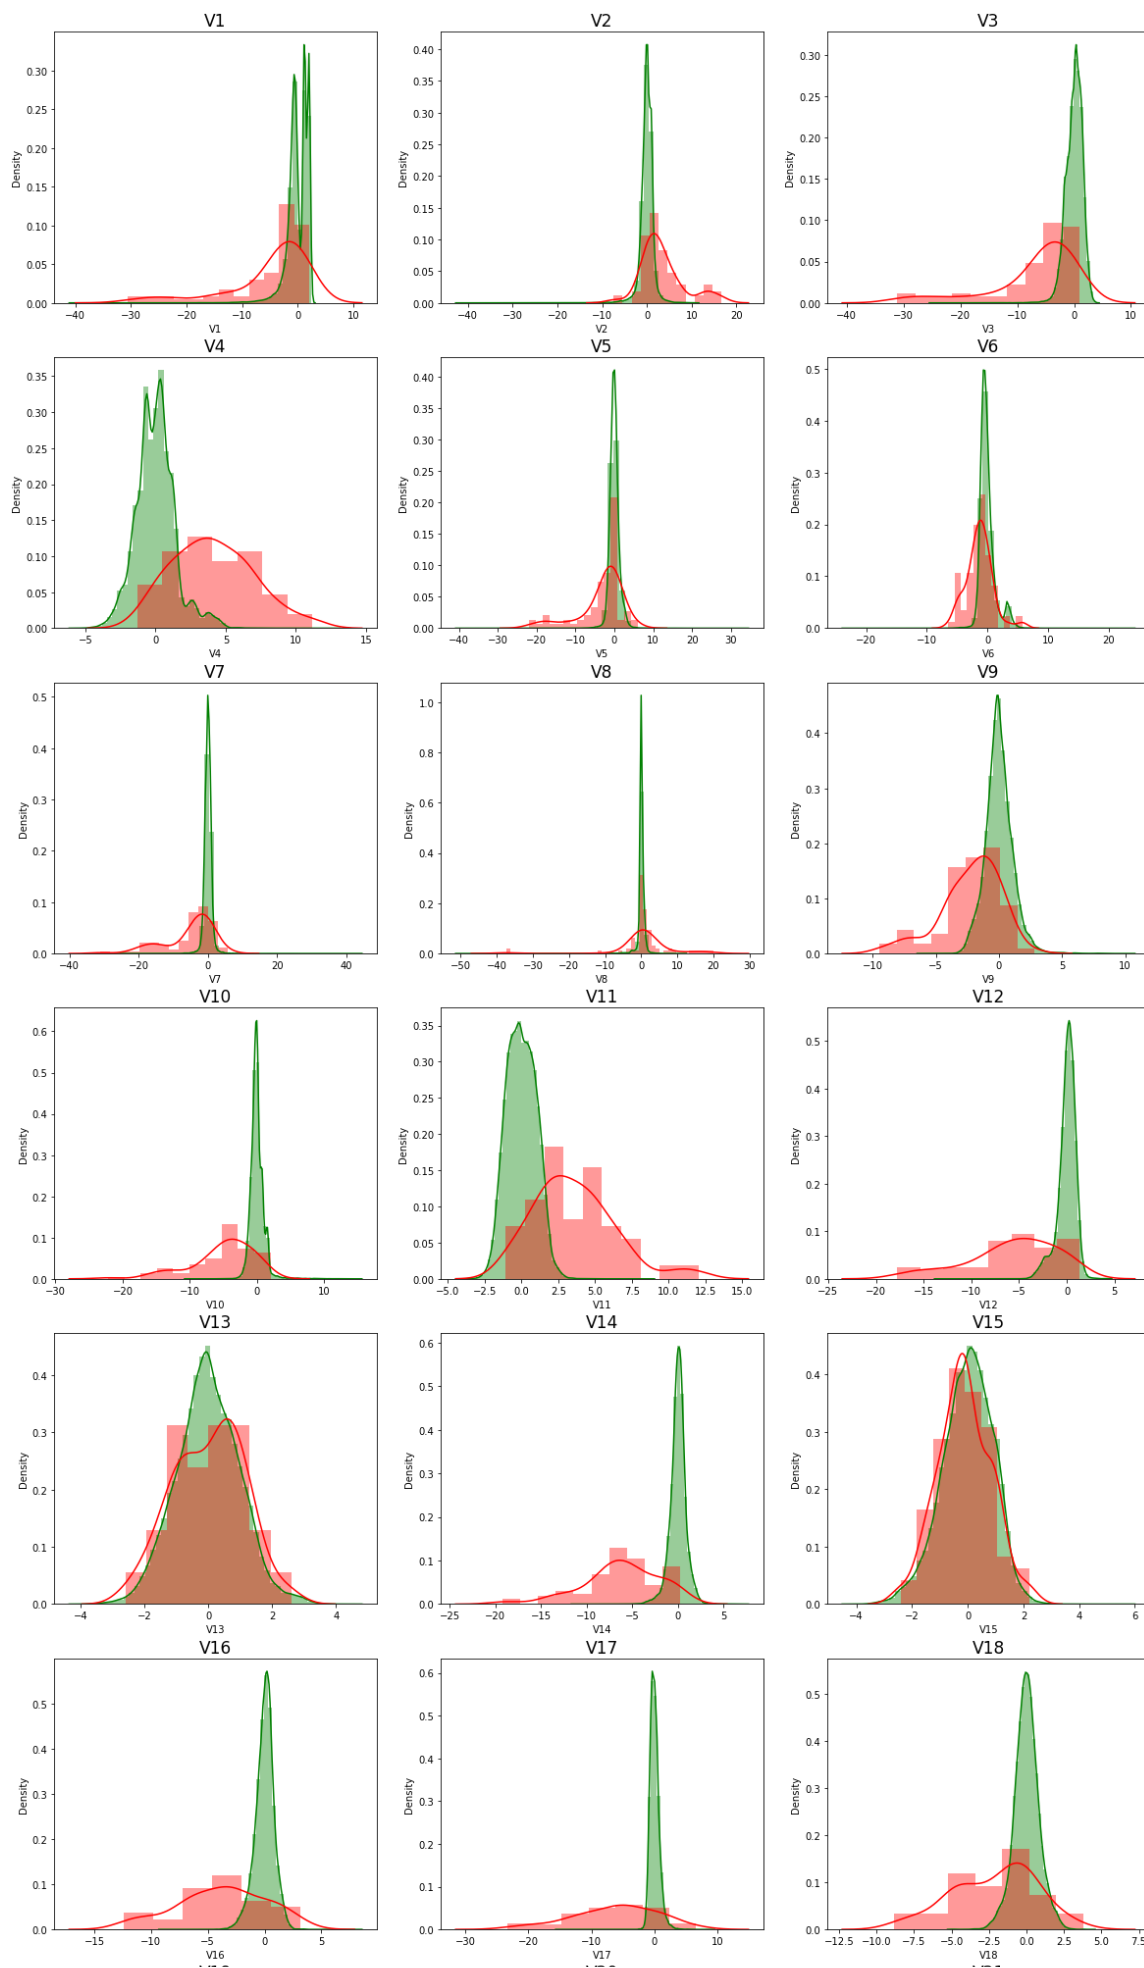
84
69
15
```

Plotting the distribution of a variable

```
In [359]: # Accumulating all the column names under one variable
cols = list(X.columns.values)
```

```
In [360]: # plot the histogram of a variable from the dataset to see the skewness
normal_records = df.Class == 0
fraud_records = df.Class == 1

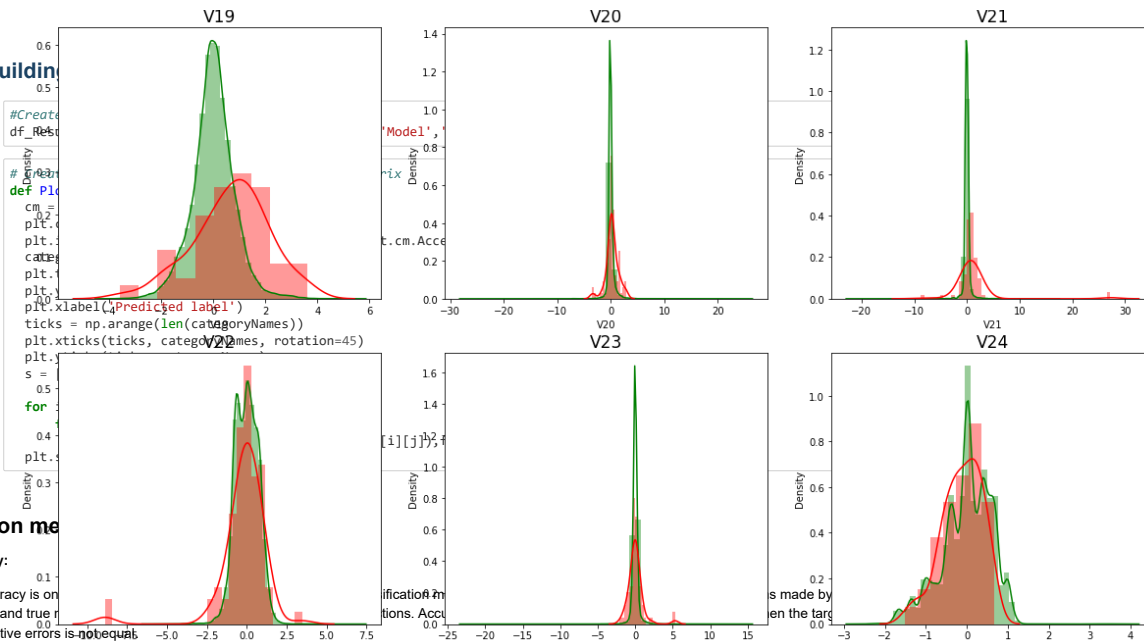
plt.figure(figsize=(20, 60))
for n, col in enumerate(cols):
    plt.subplot(10, 3, n+1)
    sns.distplot(X[col][normal_records], color='green')
    sns.distplot(X[col][fraud_records], color='red')
    plt.title(col, fontsize=17)
plt.show()
```

Model Building

```
In [361]: #Create
df_Res

In [362]: # Create
def Pl
```



Evaluation me

1. Accuracy:

- Accuracy is on (TP) and true negative errors is not equal.

2. ROC-AUC Score:

- The receiver operating rate (1)

3. Confusion Matrix:

- The confusion values for true score, among
 - False Neg
 - True Positi
 - False Pos
 - True negat
 - True Posi
 - True Neg

class = negative and predicted class = negative.

ification

ations. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

tions. Acco

Logistic Regre

- Logistic regression
- The logistic regress then be transforme
- L1 Regularization**
 - Adds a penalty
 - Leads to spars
 - Result in featur
- L2 Regularization**
 - Adds a penalty
 - Leads to soluti
 - Result in shrinki

Cross Validation:

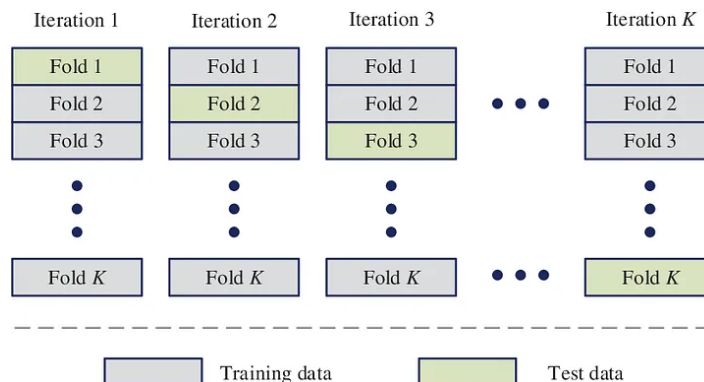

- Cross Validation is a technique for evaluating the performance of machine learning models.

- The idea is to split the available data into several parts (folds), use one part for testing and the remaining parts for training, repeat this process multiple times with different folds being used as the test set, and finally average the results to get a better estimate of the model's performance.

K-Fold Cross Validation:

- K-Fold Cross Validation is a specific implementation of cross validation where the data is divided into k equal parts, or folds.

- The model is trained on k-1 folds and tested on the remaining one, the process is repeated k times with each fold being used as the test set once. The average performance across all k iterations is used as the performance measure for the model.



Training data Test data

```

In [363]: ## Created a common function to fit and predict on a Logistic Regression model for both L1 and L2
def buildAndRunLogisticModels(df_Results, Methodology, X_train,y_train, X_test, y_test ):

    # Logistic Regression
    from sklearn import linear_model
    from sklearn.model_selection import KFold

    num_C = list(np.power(10,0, np.arange(-10, 10)))
    cv_num = KFold(n_splits=10, shuffle=True, random_state=42)
    # CV is used for cross validation in LogisticRegressionCV
    searchCV_l2 = linear_model.LogisticRegressionCV(
        Cs= num_C
        ,penalty='l2'
        ,scoring='roc_auc'
        ,cv=cv_num
        ,random_state=42
        ,max_iter=10000
        ,fit_intercept=True
        ,solver='newton-cg'
        ,tol=10
    )

    searchCV_l1 = linear_model.LogisticRegressionCV(
        Cs=num_C
        ,penalty='l1'
        ,scoring='roc_auc'
        ,cv=cv_num
        ,random_state=42
        ,max_iter=10000
        ,fit_intercept=True
        ,solver='liblinear'
        ,tol=10
    )

    searchCV_l1.fit(X_train, y_train)
    searchCV_l2.fit(X_train, y_train)
    print ('Max auc_roc for l1:', searchCV_l1.scores_[1].mean(axis=0).max())
    print ('Max auc_roc for l2:', searchCV_l2.scores_[1].mean(axis=0).max())

    print("Parameters for l1 regularisations")
    print(searchCV_l1.coef_)
    print(searchCV_l1.intercept_)
    print(searchCV_l1.scores_)

    print("Parameters for l2 regularisations")
    print(searchCV_l2.coef_)
    print(searchCV_l2.intercept_)
    print(searchCV_l2.scores_)

    #find predicted values
    y_pred_l1 = searchCV_l1.predict(X_test)
    y_pred_l2 = searchCV_l2.predict(X_test)

    #Find predicted probabilities
    y_pred_probs_l1 = searchCV_l1.predict_proba(X_test)[:,:1]
    y_pred_probs_l2 = searchCV_l2.predict_proba(X_test)[:,:1]

    # Accuracy of L2/L1 models
    Accuracy_l2 = metrics.accuracy_score(y_pred=y_pred_l2, y_true=y_test)
    Accuracy_l1 = metrics.accuracy_score(y_pred=y_pred_l1, y_true=y_test)

    print("Accuracy of Logistic model with l2 regularisation : {}".format(Accuracy_l2))
    print("Confusion Matrix")
    Plot_confusion_matrix(y_test, y_pred_l2)
    print("classification Report")
    print(classification_report(y_test, y_pred_l2))

    print("Accuracy of Logistic model with l1 regularisation : {}".format(Accuracy_l1))
    print("Confusion Matrix")
    Plot_confusion_matrix(y_test, y_pred_l1)
    print("classification Report")
    print(classification_report(y_test, y_pred_l1))

    l2_roc_value = roc_auc_score(y_test, y_pred_probs_l2)
    print("l2 roc_value: {}".format(l2_roc_value))
    fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred_probs_l2)
    threshold = thresholds[np.argmax(tpr-fpr)]
    print("l2 threshold: {}".format(threshold))

    roc_auc = metrics.auc(fpr, tpr)
    print("ROC for the test dataset", '{:.1%}'.format(roc_auc))
    plt.plot(fpr,tpr,label="Test, auc="+str(roc_auc))
    plt.legend(loc=4)
    plt.show()

    df_Results = df_Results.append(pd.DataFrame({'Methodology': Methodology, 'Model': 'Logistic Regression with L2 Regularisation', 'Accuracy': Accuracy_l2, 'roc_value': l2_roc_value, 'threshold': threshold}, index=[0]), ignore_index= True)

    l1_roc_value = roc_auc_score(y_test, y_pred_probs_l1)
    print("l1 roc_value: {}".format(l1_roc_value))
    fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred_probs_l1)
    threshold = thresholds[np.argmax(tpr-fpr)]
    print("l1 threshold: {}".format(threshold))

    roc_auc = metrics.auc(fpr, tpr)
    print("ROC for the test dataset", '{:.1%}'.format(roc_auc))
    plt.plot(fpr,tpr,label="Test, auc="+str(roc_auc))
    plt.legend(loc=4)
    plt.show()

    df_Results = df_Results.append(pd.DataFrame({'Methodology': Methodology, 'Model': 'Logistic Regression with L1 Regularisation', 'Accuracy': Accuracy_l1, 'roc_value': l1_roc_value, 'threshold': threshold}, index=[0]), ignore_index= True)
    return df_Results

```

KNeighborsClassifier

- KNeighborsClassifier is a class in the scikit-learn library that implements the K-Nearest Neighbors (KNN) algorithm for classification.
- It is a simple and effective algorithm for solving classification problems by finding the closest training examples and assigning the majority class of these nearest neighbors to the test sample. The number of nearest neighbors (k) and the distance metric used can be specified as parameters during model training.

```
In [364]: # Created a common function to fit and predict on a KNN model
def buildAndRunKNNModels(df_Results,Methodology, X_train,y_train, X_test, y_test ):

    #create KNN model and fit the model with train dataset
    knn = KNeighborsClassifier(n_neighbors = 5,n_jobs=16)
    knn.fit(X_train,y_train)
    score = knn.score(X_test,y_test)
    print("model score")
    print(score)

    #Accuracy
    y_pred = knn.predict(X_test)
    KNN_Accuracy = metrics.accuracy_score(y_pred=y_pred, y_true=y_test)
    print("Confusion Matrix")
    Plot_confusion_matrix(y_test, y_pred)
    print("classification Report")
    print(classification_report(y_test, y_pred))

    knn_probs = knn.predict_proba(X_test)[: , 1]

    # Calculate roc auc
    knn_roc_value = roc_auc_score(y_test, knn_probs)
    print("KNN roc_value: {0}" .format(knn_roc_value))
    fpr, tpr, thresholds = metrics.roc_curve(y_test, knn_probs)
    threshold = thresholds[np.argmax(tpr-fpr)]
    print("KNN threshold: {0}" .format(threshold))

    roc_auc = metrics.auc(fpr, tpr)
    print("ROC for the test dataset", '{:.1%}'.format(roc_auc))
    plt.plot(fpr,tpr,label="Test, auc="+str(roc_auc))
    plt.legend(loc=4)
    plt.show()

    df_Results = df_Results.append(pd.DataFrame({'Methodology': Methodology,'Model': 'KNN','Accuracy': score,'roc_value': knn_roc_value,'threshold': threshold}, index=[0]),ignore_index= True)

    return df_Results
```

Decision Tree Classifier

- Predictive Model:
 - Decision Tree Classifier is a supervised learning algorithm used for classification problems. It creates a tree-like model of decisions and their possible consequences, which can be used to predict the class of a new data point. The algorithm starts at the root node and splits the data into different branches based on the values of the features. It continues this process recursively until it reaches the leaf nodes, which represent the final prediction.
- Tree Representation:
 - Decision trees can be easily visualized, which makes them a useful tool for interpreting the relationship between the features and the target variable. The tree representation also makes it possible to understand the logic behind the model's predictions. Each internal node in the tree represents a feature, and each branch represents a possible value of that feature. The leaves of the tree represent the class predictions for the instances that reach that leaf.

```
In [365]: # Created a common function to fit and predict on a Tree models for both gini and entropy criteria
def buildAndRunTreeModels(df_Results, Methodology, X_train,y_train, X_test, y_test ):
    #Evaluate Decision Tree model with 'gini' & 'entropy'
    criteria = ['gini', 'entropy']
    scores = {}

    for c in criteria:
        dt = DecisionTreeClassifier(criterion = c, random_state=42)
        dt.fit(X_train, y_train)
        y_pred = dt.predict(X_test)
        test_score = dt.score(X_test, y_test)
        tree_preds = dt.predict_proba(X_test)[: , 1]
        tree_roc_value = roc_auc_score(y_test, tree_preds)
        scores = test_score
        print(c + " score: {0}" .format(test_score))
        print("Confusion Matrix")
        Plot_confusion_matrix(y_test, y_pred)
        print("classification Report")
        print(classification_report(y_test, y_pred))
        print(c + " tree_roc_value: {0}" .format(tree_roc_value))
        fpr, tpr, thresholds = metrics.roc_curve(y_test, tree_preds)
        threshold = thresholds[np.argmax(tpr-fpr)]
        print("Tree threshold: {0}" .format(threshold))
        roc_auc = metrics.auc(fpr, tpr)
        print("ROC for the test dataset", '{:.1%}'.format(roc_auc))
        plt.plot(fpr,tpr,label="Test, auc="+str(roc_auc))
        plt.legend(loc=4)
        plt.show()

        df_Results = df_Results.append(pd.DataFrame({'Methodology': Methodology,'Model': 'Tree Model with {0} criteria'.format(c),'Accuracy': test_score,'roc_value': tree_roc_value,'threshold': threshold}, index=[0]),ignore_index= True)

    return df_Results
```

Random Forest Classifier

- Ensemble Method:
 - Random Forest Classifier is an ensemble learning method for classification problems. It creates multiple decision trees and combines their predictions to produce a final output. This helps to reduce overfitting, which is a common issue with decision trees. The idea behind this method is to train multiple trees on random subsets of the data, and average their predictions to produce a more robust result.
- Bagging and Feature Selection:
 - Random Forest Classifier is based on two main concepts: Bagging (Bootstrapped Aggregating) and feature selection. Bagging is a resampling technique used to create multiple training sets from the original data. In Random Forest Classifier, each decision tree is trained on a different bootstrapped sample of the data. Feature selection is used to randomly select a subset of features for each split, which helps to reduce the correlation between trees and improve the overall accuracy of the model. The combination of bagging and feature selection results in a strong and reliable prediction model.

```
In [366]: # Created a common function to fit and predict on a Random Forest model
def buildAndRunRandomForestModels(df_Results, Methodology, X_train,y_train, X_test, y_test ):
    #Evaluate Random Forest model

    # Create the model with 100 trees
    RF_model = RandomForestClassifier(n_estimators=100,
                                     bootstrap = True,
                                     max_features = 'sqrt', random_state=42)

    # Fit on training data
    RF_model.fit(X_train, y_train)
    RF_test_score = RF_model.score(X_test, y_test)
    RF_model.predict(X_test)

    print('Model Accuracy: {}'.format(RF_test_score))

    # Actual class predictions
    rf_predictions = RF_model.predict(X_test)

    print("Confusion Matrix")
    Plot_confusion_matrix(y_test, rf_predictions)
    print("classification Report")
    print(classification_report(y_test, rf_predictions))

    # Probabilities for each class
    rf_probs = RF_model.predict_proba(X_test)[:, 1]

    # Calculate roc auc
    roc_value = roc_auc_score(y_test, rf_probs)

    print("Random Forest roc_value: {}".format(roc_value))
    fpr, tpr, thresholds = metrics.roc_curve(y_test, rf_probs)
    threshold = thresholds[np.argmax(tpr-fpr)]
    print("Random Forest threshold: {}".format(threshold))
    roc_auc = metrics.auc(fpr, tpr)
    print("ROC for the test dataset", '{:.1%}'.format(roc_auc))
    plt.plot(fpr,tpr,label="Test, auc="+str(roc_auc))
    plt.legend(loc=4)
    plt.show()

    df_Results = df_Results.append(pd.DataFrame({'Methodology': Methodology, 'Model': 'Random Forest', 'Accuracy': RF_test_score, 'roc_value': roc_value, 'threshold': threshold}, index=[0]), ignore_index= True)

    return df_Results
```

XGBoost (eXtreme Gradient Boosting) Classifier

- Gradient Boosting:
 - XGBoost (eXtreme Gradient Boosting) Classifier is an implementation of gradient boosting for classification problems. It is an advanced version of gradient boosting that is designed to handle large datasets and high-dimensional features. XGBoost Classifier trains weak decision trees in a sequential manner, using the errors from the previous tree to improve the predictions of the next tree. This process continues until a set number of trees have been trained, or a specified stopping criterion is met.
- Scalability and Performance:
 - XGBoost Classifier is known for its scalability and performance. It uses parallel processing and efficient algorithms to train models quickly and accurately. It also includes a number of optimization techniques, such as regularization, to help prevent overfitting and improve the generalization performance of the model. XGBoost Classifier has won several machine learning competitions and is widely used in industry due to its ability to handle large datasets and produce accurate predictions.

```
In [367]: # Created a common function to fit and predict on a XGBoost model
def buildAndRunXGBoostModels(df_Results, Methodology,X_train,y_train, X_test, y_test ):
    #Evaluate XGboost model
    XGBmodel = XGBClassifier(random_state=42)
    XGBmodel.fit(X_train, y_train)
    y_pred = XGBmodel.predict(X_test)

    XGB_test_score = XGBmodel.score(X_test, y_test)
    print('Model Accuracy: {}'.format(XGB_test_score))

    print("Confusion Matrix")
    Plot_confusion_matrix(y_test, y_pred)
    print("classification Report")
    print(classification_report(y_test, y_pred))
    # Probabilities for each class
    XGB_probs = XGBmodel.predict_proba(X_test)[:, 1]

    # Calculate roc auc
    XGB_roc_value = roc_auc_score(y_test, XGB_probs)

    print("XGboost roc_value: {}".format(XGB_roc_value))
    fpr, tpr, thresholds = metrics.roc_curve(y_test, XGB_probs)
    threshold = thresholds[np.argmax(tpr-fpr)]
    print("XGBoost threshold: {}".format(threshold))
    roc_auc = metrics.auc(fpr, tpr)
    print("ROC for the test dataset", '{:.1%}'.format(roc_auc))
    plt.plot(fpr,tpr,label="Test, auc="+str(roc_auc))
    plt.legend(loc=4)
    plt.show()

    df_Results = df_Results.append(pd.DataFrame({'Methodology': Methodology, 'Model': 'XGBoost', 'Accuracy': XGB_test_score, 'roc_value': XGB_roc_value, 'threshold': threshold}, index=[0]), ignore_index= True)

    return df_Results
```

Support Vector Classifier (SVC)

- Support Vector Machines:
 - Support Vector Classifier (SVC) is a type of Support Vector Machine (SVM) algorithm used for classification problems. SVC is a linear classifier that finds the hyperplane that best separates the data into classes. The goal of the algorithm is to maximize the margin between the hyperplane and the closest data points, called support vectors, which are used to define the boundary between the classes.
- Kernel Trick:
 - SVC can also be used for non-linearly separable data by using a technique called the kernel trick. The kernel trick transforms the data into a high-dimensional feature space, where a linear boundary can be found. The transformation is performed implicitly, so the user does not need to explicitly perform the calculation. This makes SVC a versatile algorithm that can be used for a wide range of classification problems, including both linearly separable and non-linearly separable data.
- Note: SVC can also be used for regression problems, in which case it is called Support Vector Regression (SVR).

```
In [368]: # Created a common function to fit and predict on a SVM model
def buildAndRunSVMModels(df_Results, Methodology, X_train,y_train, X_test, y_test ):
    #Evaluate SVM model with sigmoid kernel model
    from sklearn.svm import SVC
    from sklearn.metrics import accuracy_score
    from sklearn.metrics import roc_auc_score

    clf = SVC(kernel='sigmoid', random_state=42)
    clf.fit(X_train,y_train)
    y_pred_SVM = clf.predict(X_test)
    SVM_Score = accuracy_score(y_test,y_pred_SVM)
    print("accuracy_score : {}".format(SVM_Score))
    print("Confusion Matrix")
    Plot_confusion_matrix(y_test, y_pred_SVM)
    print("classification Report")
    print(classification_report(y_test, y_pred_SVM))

    # Run classifier
    classifier = SVC(kernel='sigmoid' , probability=True)
    svm_probs = classifier.fit(X_train, y_train).predict_proba(X_test)[:, 1]

    # Calculate roc auc
    roc_value = roc_auc_score(y_test, svm_probs)

    print("SVM roc_value: {}".format(roc_value))
    fpr, tpr, thresholds = metrics.roc_curve(y_test, svm_probs)
    threshold = thresholds[np.argmax(tpr-fpr)]
    print("SVM threshold: {}".format(threshold))
    roc_auc = metrics.auc(fpr, tpr)
    print("ROC for the test dataset", '{:.1%}'.format(roc_auc))
    plt.plot(fpr,tpr,label="Test, auc="+str(roc_auc))
    plt.legend(loc=4)
    plt.show()

    df_Results = df_Results.append(pd.DataFrame({'Methodology': Methodology, 'Model': 'SVM', 'Accuracy': SVM_Score, 'roc_value': roc_value, 'threshold': threshold}, index=[0]), ignore_index= True)

    return df_Results
```

ANN

```
In [369]: # Created a common function to fit and predict on a SVM model
def buildAndgetmetricsANN(df_Results, Methodology, X_train,y_train, X_test, y_test ):
    from sklearn.metrics import accuracy_score
    from sklearn.metrics import roc_auc_score

    ann = Sequential([Dense(input_dim = 30, units = 16, activation = 'relu'),
        Dense(units = 24, activation = 'relu'),
        Dropout(0.5),
        Dense(units = 20, activation = 'relu'),
        Dense(units = 24, activation = 'relu'),
        Dense(units =1, activation = 'sigmoid'),])

    ann.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])
    ann.fit(X_train, y_train, batch_size = 15, epochs = 5)

    ann.fit(X_train,y_train)
    y_pred_ann_prob = ann.predict(X_test)
    y_pred_ann = np.where(y_pred_ann_prob>0.5,1,0)
    ANN_Score = accuracy_score(y_test,y_pred_ann)
    print("accuracy_score : {}".format(ANN_Score))
    print("Confusion Matrix")
    Plot_confusion_matrix(y_test, y_pred_ann)
    print("classification Report")
    print(classification_report(y_test, y_pred_ann))

    # Run classifier
    # classifier = SVC(kernel='sigmoid' , probability=True)
    # svm_probs = classifier.fit(X_train, y_train).predict_proba(X_test)[:, 1]

    # Calculate roc auc
    roc_value = roc_auc_score(y_test, y_pred_ann_prob)

    print("ANN roc_value: {}".format(roc_value))
    fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred_ann_prob)
    threshold = thresholds[np.argmax(tpr-fpr)]
    print("ANN threshold: {}".format(threshold))
    roc_auc = metrics.auc(fpr, tpr)
    print("ROC for the test dataset", '{:.1%}'.format(roc_auc))
    plt.plot(fpr,tpr,label="Test, auc="+str(roc_auc))
    plt.legend(loc=4)
    plt.show()

    df_Results = df_Results.append(pd.DataFrame({'Methodology': Methodology, 'Model': 'ANN', 'Accuracy': ANN_Score, 'roc_value': roc_value, 'threshold': threshold}, index=[0]), ignore_index= True)

    return df_Results
```

- Build different models on the imbalanced dataset and see the result

Perform cross validation with RepeatedKfold

RepeatedKfold

- Cross-Validation:
 - RepeatedKfold is a type of cross-validation technique used for model selection and evaluation. Cross-validation is a method for evaluating the performance of a machine learning model by dividing the data into training and testing sets, and training the model on the training set and evaluating it on the testing set. RepeatedKfold is a variation of k-fold cross-validation, where the same k-folds are used multiple times to get a better estimate of the model's performance.
- Estimating Model Performance:
 - The purpose of RepeatedKfold is to get a more reliable estimate of the model's performance by repeating the k-fold process multiple times and averaging the results. By repeating the process, RepeatedKfold helps to reduce the variance of the estimate, making it a more robust evaluation technique. Additionally, by using the same folds multiple times, the user can ensure that all parts of the data are used for both training and testing, giving a more comprehensive evaluation of the model's performance.

In [370]:

```
#Lets perfrom RepeatedKFold and check the results
from sklearn.model_selection import RepeatedKFold
rkf = RepeatedKFold(n_splits=5, n_repeats=10, random_state=None)
# X is the feature set and y is the target
for train_index, test_index in rkf.split(X,y):
    print("TRAIN:", train_index, "TEST:", test_index)
    X_train_cv, X_test_cv = X.iloc[train_index], X.iloc[test_index]
    y_train_cv, y_test_cv = y.iloc[train_index], y.iloc[test_index]
```

```
TRAIN: [ 0  2  3 ... 49994 49996 49999] TEST: [  1  8 22 ... 49995 49997 49998]
TRAIN: [ 0  1  2 ... 49997 49998 49999] TEST: [ 13 16 18 ... 49989 49994 49996]
TRAIN: [ 1  2  3 ... 49996 49997 49998] TEST: [  0  4  7 ... 49992 49993 49999]
TRAIN: [ 0  1  2 ... 49997 49998 49999] TEST: [  5  9 14 ... 49963 49977 49978]
TRAIN: [ 0  1  4 ... 49997 49998 49999] TEST: [  2  3  6 ... 49975 49987 49990]
TRAIN: [ 0  1  3 ... 49996 49997 49999] TEST: [  2 11 55 ... 49983 49991 49998]
TRAIN: [ 0  1  2 ... 49994 49995 49998] TEST: [  4  5 10 ... 49996 49997 49999]
TRAIN: [ 0  2  4 ... 49997 49998 49999] TEST: [  1  3 14 ... 49984 49993 49995]
TRAIN: [ 1  2  3 ... 49997 49998 49999] TEST: [  0  7  9 ... 49988 49990 49994]
TRAIN: [ 0  1  2 ... 49997 49998 49999] TEST: [  6  8 16 ... 49978 49987 49989]
TRAIN: [ 0  1  2 ... 49994 49996 49997] TEST: [  6  9 10 ... 49995 49998 49999]
TRAIN: [ 0  4  5 ... 49996 49998 49999] TEST: [  1  2  3 ... 49984 49987 49997]
TRAIN: [ 1  2  3 ... 49997 49998 49999] TEST: [  0  5 11 ... 49990 49991 49996]
TRAIN: [ 0  1  2 ... 49997 49998 49999] TEST: [  7  8 13 ... 49975 49977 49993]
TRAIN: [ 0  1  2 ... 49997 49998 49999] TEST: [  4 19 25 ... 49982 49989 49994]
TRAIN: [ 0  1  2 ... 49997 49998 49999] TEST: [  8 21 24 ... 49984 49993 49994]
TRAIN: [ 0  2  5 ... 49997 49998 49999] TEST: [  1  3  4 ... 49975 49982 49987]
TRAIN: [ 1  3  4 ... 49996 49997 49998] TEST: [  0  2  5 ... 49989 49995 49999]
TRAIN: [ 0  1  2 ... 49996 49998 49999] TEST: [  9 15 18 ... 49991 49992 49997]
TRAIN: [ 0  1  2 ... 49995 49997 49999] TEST: [  7 29 45 ... 49986 49996 49998]
TRAIN: [ 0  1  2 ... 49997 49998 49999] TEST: [ 11 14 19 ... 49987 49992 49996]
TRAIN: [ 0  1  2 ... 49996 49997 49998] TEST: [  4  6 12 ... 49991 49993 49999]
TRAIN: [ 1  2  4 ... 49997 49998 49999] TEST: [  0  3  7 ... 49981 49989 49994]
TRAIN: [ 0  3  4 ... 49997 49998 49999] TEST: [  1  2  5 ... 49961 49966 49990]
TRAIN: [ 0  1  2 ... 49994 49996 49999] TEST: [  8 10 22 ... 49995 49997 49998]
TRAIN: [ 0  1  2 ... 49996 49997 49999] TEST: [  4  9 14 ... 49989 49992 49998]
TRAIN: [ 0  1  2 ... 49997 49998 49999] TEST: [ 10 11 12 ... 49966 49974 49983]
TRAIN: [ 3  4  5 ... 49996 49998 49999] TEST: [  0  1  2 ... 49972 49994 49997]
TRAIN: [ 0  1  2 ... 49997 49998 49999] TEST: [  3  6  8 ... 49991 49993 49996]
TRAIN: [ 0  1  2 ... 49996 49997 49998] TEST: [  5  7 15 ... 49987 49995 49999]
TRAIN: [ 0  2  4 ... 49997 49998 49999] TEST: [  1  3  5 ... 49978 49990 49992]
TRAIN: [ 0  1  2 ... 49997 49998 49999] TEST: [  4  8 26 ... 49984 49991 49993]
TRAIN: [ 0  1  2 ... 49993 49995 49997] TEST: [  6  7 12 ... 49996 49998 49999]
TRAIN: [ 1  2  3 ... 49997 49998 49999] TEST: [  0  9 22 ... 49987 49989 49995]
TRAIN: [ 0  1  3 ... 49996 49998 49999] TEST: [  2 11 13 ... 49983 49985 49997]
TRAIN: [ 0  1  2 ... 49996 49997 49998] TEST: [  6  8 11 ... 49970 49978 49999]
TRAIN: [ 0  1  2 ... 49997 49998 49999] TEST: [  5 12 17 ... 49992 49994 49995]
TRAIN: [ 0  3  5 ... 49997 49998 49999] TEST: [  1  2  4 ... 49990 49993 49996]
TRAIN: [ 1  2  4 ... 49996 49998 49999] TEST: [  0  3  7 ... 49985 49991 49997]
TRAIN: [ 0  1  2 ... 49996 49997 49999] TEST: [  9 13 14 ... 49988 49989 49998]
TRAIN: [ 0  1  3 ... 49996 49998 49999] TEST: [  2  6 17 ... 49988 49991 49997]
TRAIN: [ 0  1  2 ... 49996 49997 49999] TEST: [  5 13 14 ... 49987 49992 49998]
TRAIN: [ 0  2  3 ... 49997 49998 49999] TEST: [  1 11 27 ... 49989 49993 49994]
TRAIN: [ 1  2  5 ... 49997 49998 49999] TEST: [  0  3  4 ... 49982 49995 49996]
TRAIN: [ 0  1  2 ... 49996 49997 49998] TEST: [  9 15 20 ... 49979 49990 49999]
TRAIN: [ 0  1  2 ... 49994 49996 49999] TEST: [  5  8  9 ... 49995 49997 49998]
TRAIN: [ 1  2  3 ... 49996 49997 49998] TEST: [  0 13 16 ... 49981 49988 49999]
TRAIN: [ 0  1  2 ... 49997 49998 49999] TEST: [  6 14 15 ... 49992 49994 49996]
TRAIN: [ 0  2  3 ... 49997 49998 49999] TEST: [  1  4  7 ... 49985 49987 49991]
TRAIN: [ 0  1  4 ... 49997 49998 49999] TEST: [  2  3 11 ... 49984 49989 49993]
```

```

In [371]: Data_Validation_Method = "RepeatedKFold Cross Validation"
#Run Logistic Regression with L1 And L2 Regularisation
print("Logistic Regression with L1 And L2 Regularisation")
start_time = time.time()
df_Results = buildAndRunLogisticModels(df_Results, Data_Validation_Method, X_train_cv, y_train_cv, X_test_cv, y_test_cv)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print("-"*60)

#Run KNN Model
print("KNN Model")
start_time = time.time()
df_Results = buildAndRunKNNModels(df_Results, Data_Validation_Method, X_train_cv, y_train_cv, X_test_cv, y_test_cv)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print("-"*60)

#Run Decision Tree Models with 'gini' & 'entropy' criteria
print("Decision Tree Models with 'gini' & 'entropy' criteria")
start_time = time.time()
df_Results = buildAndRunTreeModels(df_Results, Data_Validation_Method, X_train_cv, y_train_cv, X_test_cv, y_test_cv)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print("-"*60)

#Run Random Forest Model
print("Random Forest Model")
start_time = time.time()
df_Results = buildAndRunRandomForestModels(df_Results, Data_Validation_Method, X_train_cv, y_train_cv, X_test_cv, y_test_cv)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print("-"*60)

#Run XGBoost Model
print("XGBoost Model")
start_time = time.time()
df_Results = buildAndRunXGBoostModels(df_Results, Data_Validation_Method, X_train_cv, y_train_cv, X_test_cv, y_test_cv)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print("-"*60)

#Run SVM Model with Sigmoid Kernel
print("SVM Model with Sigmoid Kernel")
start_time = time.time()
df_Results = buildAndRunSVMModels(df_Results, Data_Validation_Method, X_train_cv, y_train_cv, X_test_cv, y_test_cv)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))

#Run ANN Model
print("ANN Model")
start_time = time.time()
df_Results = buildAndGetMetricsANN(df_Results, Data_Validation_Method, X_train_Smote, y_train_Smote, X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print("-"*80)

```




```

Logistic Regression with L1 And L2 Regularisation
Max auc_roc for l1: 0.9632348876283263
Max auc_roc for l2: 0.9584189367560885
Parameters for l1 regularisations
[[-0.05916746 -0.14318117 -0.18150557 0.05851872 -0.27074988 0.09183386
 0.14170428 -0.03610091 -0.15757988 -0.13344612 -0.03987543 0.01458198
 -0.135869 -0.21607118 0.01094898 -0.10418107 -0.21198404 0.03788204
 0.01613639 0.30092196 0.08855907 -0.019655 0.0457783 -0.01608273
 0.02177818 0.00577888 -0.10557782 0.01302801 -0.00660377 -0.11590842]]
[-2.08620727]
{1: array([[0.5, 0.5, 0.5, 0.5, 0.41701076,
0.48761261, 0.61524024, 0.55868368, 0.99931181, 1.,
1., 1., 1., 1., 1.,
1., 0.99993744, 1., 1., 1., 1.,
0.5, 0.5, 0.5, 0.5, 0.59346029,
0.59134441, 0.57185612, 0.58018041, 0.87407779, 0.87482948,
0.84208915, 0.86185584, 0.86670008, 0.86152176, 0.85394916,
0.84520727, 0.84078065, 0.87352098, 0.86344275, 0.90253069],
0.5, 0.5, 0.5, 0.5, 0.45969925,
0.43964912, 0.49819549, 0.5385213, 0.87102757, 0.93265664,
0.82075188, 0.88501253, 0.93275689, 0.91591479, 0.83488722,
0.93857143, 0.8712782, 0.9141604, 0.86055138, 0.99308271],
0.5, 0.5, 0.5, 0.5, 0.58969925,
0.56963659, 0.60986216, 0.59397243, 0.77929825, 0.86140351,
0.83300752, 0.83729323, 0.88779449, 0.85674185, 0.89077694,
0.89263158, 0.82255639, 0.86701754, 0.91468672, 0.96704261],
0.5, 0.5, 0.5, 0.5, 0.66948198,
0.62831582, 0.5640015, 0.58427177, 0.75337838, 0.81637888,
0.68512262, 0.71458959, 0.81206206, 0.74399399, 0.81781782,
0.86286286, 0.69901151, 0.77752753, 0.86273774, 0.85635636],
0.5, 0.5, 0.5, 0.5, 0.39401095,
0.43071804, 0.44542235, 0.42191693, 0.62448571, 0.88089872,
0.85052413, 0.86633752, 0.89817896, 0.8750313, 0.905549,
0.94157633, 0.87746413, 0.89495904, 0.90744517, 0.98615434],
0.5, 0.5, 0.5, 0.5, 0.29233301,
0.3078423, 0.36955744, 0.3192909, 0.84998748, 0.98415084,
0.88859075, 0.92290079, 0.97098494, 0.96357912, 0.90966334,
0.96060964, 0.9301635, 0.9561733, 0.69832922, 0.99062645],
0.5, 0.5, 0.5, 0.5, 0.6000167,
0.55620981, 0.53906011, 0.53279601, 0.97962081, 0.94320555,
0.91962471, 0.9282274, 0.93515966, 0.92396782, 0.91235836,
0.91227484, 0.92831092, 0.93362844, 0.91728612, 0.93758178],
0.5, 0.5, 0.5, 0.5, 0.5650025,
0.51786161, 0.47394269, 0.50203328, 0.90077578, 0.92842843,
0.80580581, 0.85003754, 0.97522523, 0.9227978, 0.9771021,
0.98936436, 0.8262012, 0.9369995, 0.99543293, 0.99912412],
0.5, 0.5, 0.5, 0.5, 0.51349186,
0.50876095, 0.4571965, 0.4852816, 0.99939925, 0.99964956,
0.98493116, 0.98763454, 0.99964956, 0.99564456, 0.99964956,
0.99984981, 0.98683354, 0.99964956, 0.99984981, 0.99984981]]})

Parameters for l2 regularisations
[[ 5.99668738e-03 -3.01823363e-03 -1.12397706e-01 2.47784799e-01
 1.05734896e-01 -3.77486716e-02 2.10189897e-02 -1.28127106e-01
 -1.10876988e-01 -1.68084395e-01 1.60997492e-01 -1.69798369e-01
 -5.27692239e-02 -3.95079590e-01 -5.02671609e-02 -6.92493777e-02
 -8.86719602e-02 -3.14514377e-03 1.67428310e-02 7.75524119e-03
 2.16229124e-02 2.53131116e-02 2.04644575e-02 -2.78403521e-02
 4.25101475e-03 6.58098416e-03 -2.25813405e-03 -2.10530722e-02
 2.82424821e-04 -4.49599848e-02]]
[-6.57797975]
{1: array([[0.6788038, 0.67892893, 0.68218218, 0.7078954, 0.83383383,
0.97378629, 0.99937437, 0.99993744, 0.99981231, 0.99968719,
0.99968719, 0.99968719, 0.99968719, 0.99968719, 0.99968719,
0.99968719, 0.99968719, 0.99968719, 0.99968719, 0.99968719],
0.53590022, 0.5360951, 0.53381219, 0.57281662, 0.74812773,
0.8629973, 0.88268048, 0.88281968, 0.88680086, 0.91695203,
0.91695203, 0.91695203, 0.91695203, 0.91695203, 0.91695203,
0.91695203, 0.91695203, 0.91695203, 0.91695203, 0.91695203],
0.6664411, 0.66641604, 0.66934837, 0.6935589, 0.7981203,
0.88799499, 0.92042607, 0.94433584, 0.97922306, 0.97385965,
0.97385965, 0.97385965, 0.97385965, 0.97385965, 0.97385965,
0.97385965, 0.97385965, 0.97385965, 0.97385965, 0.97385965],
0.4497995, 0.4497995, 0.45072682, 0.46401003, 0.57325815,
0.72536341, 0.84661654, 0.91932331, 0.94325815, 0.94285714,
0.94285714, 0.94285714, 0.94285714, 0.94285714, 0.94285714,
0.94285714, 0.94285714, 0.94285714, 0.94285714, 0.94285714],
0.32995495, 0.32995495, 0.33051802, 0.33508509, 0.37875375,
0.52195946, 0.6971972, 0.82388639, 0.84922422, 0.80061311,
0.80061311, 0.80061311, 0.80061311, 0.80061311, 0.80061311,
0.80061311, 0.80061311, 0.80061311, 0.80061311, 0.80061311],
0.67936746, 0.67936746, 0.68158563, 0.70086938, 0.80823584,
0.92919752, 0.94633466, 0.95985832, 0.97928518, 0.96865944,
0.96865944, 0.96865944, 0.96865944, 0.96865944, 0.96865944,
0.96865944, 0.96865944, 0.96865944, 0.96865944, 0.96865944],
0.84347608, 0.84347608, 0.84540804, 0.86519266, 0.94232764,
0.98103825, 0.98958892, 0.99556366, 0.99445458, 0.99484813,
0.99484813, 0.99484813, 0.99484813, 0.99484813, 0.99484813,
0.99484813, 0.99484813, 0.99484813, 0.99484813, 0.99484813],
0.50034801, 0.50034801, 0.50232467, 0.5316128, 0.70564325,
0.89362176, 0.9471032, 0.95448092, 0.95314458, 0.92931318,
0.92931318, 0.92931318, 0.92931318, 0.92931318, 0.92931318,
0.92931318, 0.92931318, 0.92931318, 0.92931318, 0.92931318],
0.45883383, 0.45883383, 0.4608984, 0.47472472, 0.57157157,
0.76332583, 0.95245245, 0.99843594, 0.99918669, 0.99931181,
0.99931181, 0.99931181, 0.99931181, 0.99931181, 0.99931181,
0.99931181, 0.99931181, 0.99931181, 0.99931181, 0.99931181],
0.59389237, 0.59389237, 0.59549437, 0.62187735, 0.78478098,
0.96290363, 0.999199, 0.99974969, 0.99979975, 0.99984981,
0.99984981, 0.99984981, 0.99984981, 0.99984981, 0.99984981,
0.99984981, 0.99984981, 0.99984981, 0.99984981, 0.99984981]]})
Accuracy of Logistic model with l2 regularisation : 0.999
Confusion Matrix


```



Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9985 |
| 1 | 0.86 | 0.40 | 0.55 | 15 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.93 | 0.70 | 0.77 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |


Accuracy of Logistic model with l1 regularisation : 0.9991
Confusion Matrix




Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9985 |
| 1 | 0.88 | 0.47 | 0.61 | 15 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.94 | 0.73 | 0.80 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

12 roc_value: 0.9508863294942415
12 threshold: 0.030689645354979082
ROC for the test dataset 95.1%




11 roc_value: 0.9335803705558338
11 threshold: 0.05783283302624465
ROC for the test dataset 93.4%



Time Taken by Model: --- 51.19481015205383 seconds ---


KNN Model
model score
0.999
Confusion Matrix



Classification Report


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9985 |
| 1 | 0.86 | 0.40 | 0.55 | 15 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.93 | 0.70 | 0.77 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

KNN roc_value: 0.8662961108329161
 KNN threshold: 0.2
 ROC for the test dataset 86.6%



Time Taken by Model: --- 45.10881972312927 seconds ---


Decision Tree Models with 'gini' & 'entropy' criteria
 gini score: 0.9989
 Confusion Matrix




Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9985 |
| 1 | 0.62 | 0.67 | 0.65 | 15 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.81 | 0.83 | 0.82 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

gini tree_roc_value: 0.8330328826573192
 Tree threshold: 1.0
 ROC for the test dataset 83.3%




entropy score: 0.9992
 Confusion Matrix



Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9985 |
| 1 | 0.71 | 0.80 | 0.75 | 15 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.85 | 0.90 | 0.87 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

entropy tree_roc_value: 0.899749624436655
Tree threshold: 1.0
ROC for the test dataset 90.0%




Time Taken by Model: --- 3.956048011779785 seconds ---

Random Forest Model

Model Accuracy: 0.9997


Confusion Matrix



Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9985 |
| 1 | 0.93 | 0.87 | 0.90 | 15 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.96 | 0.93 | 0.95 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

Random Forest roc_value: 0.9653814054414955
Random Forest threshold: 0.3
ROC for the test dataset 96.5%




Time Taken by Model: --- 22.89270567893982 seconds ---

XGBoost Model

Model Accuracy: 0.9996


Confusion Matrix



Classification Report


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9985 |
| 1 | 0.87 | 0.87 | 0.87 | 15 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.93 | 0.93 | 0.93 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

XGboost roc_value: 0.9363645468202302
 XGBoost threshold: 0.28312182426452637
 ROC for the test dataset 93.6%



Time Taken by Model: --- 9.490098476409912 seconds ---


SVM Model with Sigmoid Kernel
 accuracy_score : 0.9981
 Confusion Matrix



Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9985 |
| 1 | 0.00 | 0.00 | 0.00 | 15 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.50 | 0.50 | 0.50 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

SVM roc_value: 0.7244867300951428
 SVM threshold: 0.0034816380928626425
 ROC for the test dataset 72.4%



Time Taken by Model: --- 6.716370582580566 seconds ---

ANN Model

Epoch 1/5
2130/2130 [=====] - 8s 3ms/step - loss: 0.1777 - accuracy: 0.9484

Epoch 2/5
2130/2130 [=====] - 5s 3ms/step - loss: 0.0412 - accuracy: 0.9887

Epoch 3/5
2130/2130 [=====] - 7s 3ms/step - loss: 0.0343 - accuracy: 0.9907

Epoch 4/5
2130/2130 [=====] - 6s 3ms/step - loss: 0.0242 - accuracy: 0.9944


Epoch 5/5
2130/2130 [=====] - 6s 3ms/step - loss: 0.0181 - accuracy: 0.9956

999/999 [=====] - 3s 3ms/step - loss: 0.0138 - accuracy: 0.9967

313/313 [=====] - 1s 2ms/step

accuracy_score : 0.9984

Confusion Matrix




Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9985 |
| 1 | 0.48 | 0.73 | 0.58 | 15 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.74 | 0.87 | 0.79 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

ANN roc_value: 0.9272775830412285

ANN threshold: 0.000521167356052299

ROC for the test dataset 92.7%



Time Taken by Model: --- 38.665369749069214 seconds ---

In [372]: # Checking the df_result dataframe which contains consolidated results of all the runs

df_Results

Out[372]:

| | Methodology | Model | Accuracy | roc_value | threshold |
|---|--------------------------------|--|----------|-----------|-----------|
| 0 | RepeatedKFold Cross Validation | Logistic Regression with L2 Regularisation | 0.9990 | 0.950886 | 0.030690 |
| 1 | RepeatedKFold Cross Validation | Logistic Regression with L1 Regularisation | 0.9991 | 0.933580 | 0.057833 |
| 2 | RepeatedKFold Cross Validation | KNN | 0.9990 | 0.866296 | 0.200000 |
| 3 | RepeatedKFold Cross Validation | Tree Model with gini criteria | 0.9989 | 0.833033 | 1.000000 |
| 4 | RepeatedKFold Cross Validation | Tree Model with entropy criteria | 0.9992 | 0.899750 | 1.000000 |
| 5 | RepeatedKFold Cross Validation | Random Forest | 0.9997 | 0.965381 | 0.300000 |
| 6 | RepeatedKFold Cross Validation | XGBoost | 0.9996 | 0.936365 | 0.283122 |
| 7 | RepeatedKFold Cross Validation | SVM | 0.9981 | 0.724487 | 0.003482 |
| 8 | RepeatedKFold Cross Validation | ANN | 0.9984 | 0.927278 | 0.000521 |


Results for cross validation with RepeatedKFold:

- Looking at Accuracy and ROC value we have "Logistic Regression with L2 Regularisation" which has provided best results for cross validation with RepeatedKFold technique

Perform cross validation with StratifiedKFold

StratifiedKFold

- StratifiedKFold is a cross-validation technique that splits a dataset into k folds while preserving the class distribution of the samples.
- This method is useful for handling imbalanced datasets, as it ensures that each fold contains a balanced representation of all class labels.



```
In [373]: #Lets perfom StratifiedKFold and check the results
from sklearn.model_selection import StratifiedKFold
skf = StratifiedKFold(n_splits=5, random_state=None)
# X is the feature set and y is the target
for train_index, test_index in skf.split(X,y):
    print("TRAIN:", train_index, "TEST:", test_index)
    X_train_SKF_cv, X_test_SKF_cv = X.iloc[train_index], X.iloc[test_index]
    y_train_SKF_cv, y_test_SKF_cv = y.iloc[train_index], y.iloc[test_index]
```

| | |
|---|---|
| TRAIN: [9053 9725 10002 ... 49997 49998 49999] | TEST: [0 1 2 ... 9999 10000 10001] |
| TRAIN: [0 1 2 ... 49997 49998 49999] | TEST: [9053 9725 10002 ... 20003 20004 20005] |
| TRAIN: [0 1 2 ... 49997 49998 49999] | TEST: [17659 18269 19032 ... 30001 30002 30003] |
| TRAIN: [0 1 2 ... 49997 49998 49999] | TEST: [29558 29634 29870 ... 40000 40001 40002] |
| TRAIN: [0 1 2 ... 40000 40001 40002] | TEST: [38516 39164 39215 ... 49997 49998 49999] |

```
In [374]: y_train_SKF_cv.value_counts()
```

```
Out[374]: 0    39933
          1     67
          Name: Class, dtype: int64
```

```
In [375]: Data_Validation_Method = "StratifiedKFold Cross Validation"
#Run Logistic Regression with L1 And L2 Regularisation
print("Logistic Regression with L1 And L2 Regularisation")
start_time = time.time()
df_Results = buildAndRunLogisticModels(df_Results,Data_Validation_Method, X_train_SKF_cv,y_train_SKF_cv, X_test_SKF_cv, y_test_SKF_cv)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*60 )

#Run KNN Model
print("KNN Model")
start_time = time.time()
df_Results = buildAndRunKNNModels(df_Results,Data_Validation_Method,X_train_SKF_cv,y_train_SKF_cv, X_test_SKF_cv, y_test_SKF_cv)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*60 )

#Run Decision Tree Models with 'gini' & 'entropy' criteria
print("Decision Tree Models with 'gini' & 'entropy' criteria")
start_time = time.time()
df_Results = buildAndRunTreeModels(df_Results,Data_Validation_Method,X_train_SKF_cv,y_train_SKF_cv, X_test_SKF_cv, y_test_SKF_cv)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*60 )

#Run Random Forest Model
print("Random Forest Model")
start_time = time.time()
df_Results = buildAndRunRandomForestModels(df_Results,Data_Validation_Method,X_train_SKF_cv,y_train_SKF_cv, X_test_SKF_cv, y_test_SKF_cv)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*60 )

#Run XGBoost Model
print("XGBoost Model")
start_time = time.time()
df_Results = buildAndRunXGBoostModels(df_Results,Data_Validation_Method,X_train_SKF_cv,y_train_SKF_cv, X_test_SKF_cv, y_test_SKF_cv)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*60 )

#Run SVM Model with Sigmoid Kernel
print("SVM Model with Sigmoid Kernel")
start_time = time.time()
df_Results = buildAndRunSVMModels(df_Results,Data_Validation_Method,X_train_SKF_cv,y_train_SKF_cv, X_test_SKF_cv, y_test_SKF_cv)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))

#Run ANN Model
print("ANN Model")
start_time = time.time()
df_Results = buildAndgetMetricsANN(df_Results, Data_Validation_Method, X_train_Smote, y_train_Smote , X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*80 )
```




```

Logistic Regression with L1 And L2 Regularisation
Max auc_roc for l1: 0.95610831614236
Max auc_roc for l2: 0.9567308642474954
Parameters for l1 regularisations
[[-0.07151745 -0.15407323 -0.19254832 0.06567634 -0.28698825 0.09795102
 0.14259043 -0.03449754 -0.15403987 -0.14443599 -0.03359353 0.00725298
 -0.12649247 -0.23510622 0.00940845 -0.11324575 -0.24276231 0.03651103
 0.0190662 0.27380617 0.08747869 -0.02202149 0.02427419 -0.00966349
 -0.00172041 0.01093092 -0.08807507 0.03092445 -0.00671194 -0.11568803]]
[-2.07376685]
{1: array([[0.5, 0.5, 0.5, 0.5, 0.39711779,
0.43189223, 0.50948622, 0.50324561, 0.77631579, 0.86656642,
0.8173183, 0.82373434, 0.870401, 0.8416792, 0.8585213,
0.86591479, 0.82308271, 0.85328321, 0.89037594, 0.95275689],
[0.5, 0.5, 0.5, 0.5, 0.52134452,
0.52622684, 0.56920798, 0.5447129, 0.94207979, 0.9265565,
0.92376064, 0.92568019, 0.91758471, 0.91387081, 0.91967117,
0.93561175, 0.92810048, 0.92029711, 0.93093807, 0.91420464],
[0.5, 0.5, 0.5, 0.5, 0.55931809,
0.54017745, 0.54447068, 0.52983793, 0.96375801, 0.98400773,
0.97255912, 0.97474151, 0.98762119, 0.98858717, 0.98926693,
0.97499195, 0.97009052, 0.98304175, 0.99116311, 0.9745984 ],
[0.5, 0.5, 0.5, 0.5, 0.49652964,
0.45758649, 0.46560052, 0.49200386, 0.8555329, 0.87907409,
0.83846732, 0.84680333, 0.87166828, 0.87986119, 0.82086509,
0.87073808, 0.86487067, 0.86741083, 0.75253122, 0.88590748],
[0.5, 0.5, 0.5, 0.5, 0.11809953,
0.11201746, 0.11995993, 0.10768845, 0.62305463, 0.87928876,
0.58473758, 0.88537083, 0.91563808, 0.90100533, 0.79120604,
0.89395728, 0.83735823, 0.91596007, 0.63568388, 0.98772852],
[0.5, 0.5, 0.5, 0.5, 0.55930611,
0.54069076, 0.57469627, 0.53245554, 0.83263402, 0.86253758,
0.82518161, 0.84030561, 0.89237851, 0.86761022, 0.89153307,
0.91223071, 0.82411698, 0.87230711, 0.90866107, 0.99120115],
[0.5, 0.5, 0.5, 0.5, 0.50397272,
0.46468343, 0.50416041, 0.5175488, 0.86161161, 0.85172673,
0.76682933, 0.82695195, 0.85441692, 0.82100851, 0.81594094,
0.82088338, 0.79416917, 0.85372873, 0.84334334, 0.89239239],
[0.5, 0.5, 0.5, 0.5, 0.65168293,
0.57929805, 0.43233859, 0.48082457, 0.88613614, 0.88951451,
0.87230981, 0.88807558, 0.92210961, 0.90102603, 0.91798048,
0.90496747, 0.87731481, 0.9115991, 0.91535285, 0.9634009 ],
[0.5, 0.5, 0.5, 0.5, 0.48923115,
0.43170191, 0.48686988, 0.44959035, 0.97323888, 0.96848056,
0.95209474, 0.95506422, 0.98229044, 0.97030518, 0.98153912,
0.98264821, 0.94941147, 0.97864119, 0.98619012, 0.99398948],
[0.5, 0.5, 0.5, 0.5, 0.83521162,
0.74782655, 0.66539659, 0.67723874, 0.99774606, 0.99678008,
0.98479482, 0.99599299, 0.99992845, 0.99685163, 0.99996422,
0.99992845, 0.99570677, 0.99978534, 0.99982112, 0.99992845]]})

Parameters for l2 regularisations
[[-1.31491217e-02 3.07463695e-02 -1.23758286e-01 2.90976907e-01
1.00763702e-01 -5.25175254e-02 4.50978185e-02 -1.31263201e-01
-1.08854534e-01 -1.56738327e-01 1.83509896e-01 -2.25403516e-01
-5.45158708e-02 -4.16368675e-01 -5.45595830e-02 -7.19437239e-02
-1.10316217e-01 -2.06331993e-02 2.04222421e-02 -2.72144122e-02
4.33597690e-02 6.29205995e-02 3.69527095e-02 -1.48824898e-02
-1.39901365e-02 2.51950981e-03 -4.00624298e-03 -7.95897195e-02
3.42612078e-04 1.20704152e-02]]
[-7.64242687]
{1: array([[0.65944862, 0.65944862, 0.66107769, 0.67919799, 0.77882206,
0.84032581, 0.88531328, 0.92263158, 0.9356391, 0.93924812,
0.93924812, 0.93924812, 0.93924812, 0.93924812, 0.93924812,
0.93924812, 0.93924812, 0.93924812, 0.93924812, 0.93924812],
[0.5307962, 0.55307962, 0.55474879, 0.58771491, 0.73147221,
0.84848105, 0.87819229, 0.90110165, 0.91103322, 0.89542647,
0.89542647, 0.89542647, 0.89542647, 0.89542647, 0.89542647,
0.89542647, 0.89542647, 0.89542647, 0.89542647, 0.89542647],
[0.49772817, 0.49772817, 0.49898036, 0.52584881, 0.67553934,
0.87138206, 0.94204143, 0.97273801, 0.98007227, 0.98443705,
0.98443705, 0.98443705, 0.98443705, 0.98443705, 0.98443705,
0.98443705, 0.98443705, 0.98443705, 0.98443705, 0.98443705],
[0.58863726, 0.58860148, 0.59024722, 0.61833208, 0.79202891,
0.9631498, 0.9797145, 0.92626382, 0.90862581, 0.91062932,
0.91062932, 0.91062932, 0.91062932, 0.91062932, 0.91062932,
0.91062932, 0.91062932, 0.91062932, 0.91062932, 0.91062932],
[0.89706987, 0.89692676, 0.89774963, 0.90143465, 0.92533362,
0.9622196, 0.93864262, 0.95302494, 0.98339952, 0.98164645,
0.98164645, 0.98164645, 0.98164645, 0.98164645, 0.98164645,
0.98164645, 0.98164645, 0.98164645, 0.98164645, 0.98164645],
[0.50973823, 0.50973823, 0.51111598, 0.53710546, 0.66846192,
0.79950526, 0.90812876, 0.97263277, 0.96940757, 0.96157941,
0.9332728, 0.9332728, 0.9332728, 0.9332728, 0.9332728,
0.9332728, 0.9332728, 0.9332728, 0.9332728, 0.9332728 ],
[0.51582833, 0.51582833, 0.51607858, 0.52152152, 0.57645145,
0.76295045, 0.90546797, 0.90934685, 0.92880038, 0.87825325,
0.87825325, 0.87825325, 0.87825325, 0.87825325, 0.87825325,
0.87825325, 0.87825325, 0.87825325, 0.87825325, 0.87825325],
[0.41166166, 0.41166166, 0.41278779, 0.44575826, 0.6265015,
0.80974725, 0.9004004, 0.94513263, 0.95952202, 0.95658158,
0.95658158, 0.94331832, 0.94331832, 0.94331832, 0.94331832,
0.94331832, 0.94331832, 0.94331832, 0.94331832, 0.94331832],
[0.554041, 0.554041, 0.55550785, 0.57096347, 0.69085185,
0.88064828, 0.9584988, 0.98533147, 0.99080534, 0.98966048,
0.98966048, 0.98966048, 0.98966048, 0.98966048, 0.98966048,
0.98966048, 0.98966048, 0.98966048, 0.98966048, 0.98966048],
[0.33032807, 0.33022074, 0.33458552, 0.38728489, 0.68140675,
0.94451004, 0.99745984, 0.99992845, 1., 1.,
1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1. ]]])
Accuracy of Logistic model with l2 regularisation : 0.9992
Confusion Matrix


```



Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.85 | 0.65 | 0.73 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.92 | 0.82 | 0.87 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |


Accuracy of Logistic model with l1 regularisation : 0.9991
Confusion Matrix




Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.79 | 0.65 | 0.71 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.89 | 0.82 | 0.85 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

l2 roc_value: 0.9566380493898451
l2 threshold: 0.002112752055838854
ROC for the test dataset 95.7%




l1 roc_value: 0.8925585259647282
l1 threshold: 0.3014627218940176
ROC for the test dataset 89.3%



Time Taken by Model: --- 54.348634481430054 seconds ---

KNN Model
model score
0.999
Confusion Matrix




Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 1.00 | 0.41 | 0.58 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 1.00 | 0.71 | 0.79 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

KNN roc_value: 0.8232495241911248

KNN threshold: 0.2

ROC for the test dataset 82.3%




Time Taken by Model: --- 44.82395100593567 seconds ---

Decision Tree Models with 'gini' & 'entropy' criteria

gini score: 0.9985

Confusion Matrix




Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.55 | 0.65 | 0.59 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.77 | 0.82 | 0.80 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

gini tree_roc_value: 0.8230786454619913


Tree threshold: 1.0

ROC for the test dataset 82.3%



entropy score: 0.999


Confusion Matrix



Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.73 | 0.65 | 0.69 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.87 | 0.82 | 0.84 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

entropy tree_roc_value: 0.8233290711857216
Tree threshold: 1.0
ROC for the test dataset 82.3%




Time Taken by Model: --- 5.102020263671875 seconds ---

Random Forest Model

Model Accuracy: 0.9995


Confusion Matrix



Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.93 | 0.76 | 0.84 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.96 | 0.88 | 0.92 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

Random Forest roc_value: 0.9371018967538934
Random Forest threshold: 0.01
ROC for the test dataset 93.7%




Time Taken by Model: --- 20.428218603134155 seconds ---

XGBoost Model

Model Accuracy: 0.9993


Confusion Matrix



Classification Report


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.92 | 0.65 | 0.76 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.96 | 0.82 | 0.88 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

XGboost roc_value: 0.9714632522346813
 XGBoost threshold: 0.0013952451990917325
 ROC for the test dataset 97.1%



Time Taken by Model: --- 9.466046333312988 seconds ---


SVM Model with Sigmoid Kernel
 accuracy_score : 0.9975
 Confusion Matrix



Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.00 | 0.00 | 0.00 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.50 | 0.50 | 0.50 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

SVM roc_value: 0.6347791245116698
 SVM threshold: 0.003424350421318629
 ROC for the test dataset 63.5%



Time Taken by Model: --- 6.3848254680633545 seconds ---

ANN Model

Epoch 1/5
 2130/2130 [=====] - 7s 3ms/step - loss: 0.1839 - accuracy: 0.9428

Epoch 2/5
 2130/2130 [=====] - 5s 2ms/step - loss: 0.0361 - accuracy: 0.9894

Epoch 3/5
 2130/2130 [=====] - 7s 3ms/step - loss: 0.0295 - accuracy: 0.9924

Epoch 4/5
 2130/2130 [=====] - 5s 2ms/step - loss: 0.0243 - accuracy: 0.9939


Epoch 5/5
 2130/2130 [=====] - 5s 2ms/step - loss: 0.0201 - accuracy: 0.9951

999/999 [=====] - 4s 4ms/step - loss: 0.0120 - accuracy: 0.9970

313/313 [=====] - 1s 2ms/step

accuracy_score : 0.9985

Confusion Matrix




Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9985 |
| 1 | 0.50 | 0.73 | 0.59 | 15 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.75 | 0.87 | 0.80 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

ANN roc_value: 0.925528292438658

ANN threshold: 7.655112858628854e-05

ROC for the test dataset 92.6%



Time Taken by Model: --- 48.55423426628113 seconds ---

In [376]: # Checking the df_result dataframe which contains consolidated results of all the runs

df_Results

Out[376]:

| | Methodology | Model | Accuracy | roc_value | threshold |
|----|----------------------------------|--|----------|-----------|-----------|
| 0 | RepeatedKFold Cross Validation | Logistic Regression with L2 Regularisation | 0.9990 | 0.950886 | 0.030690 |
| 1 | RepeatedKFold Cross Validation | Logistic Regression with L1 Regularisation | 0.9991 | 0.933580 | 0.057833 |
| 2 | RepeatedKFold Cross Validation | KNN | 0.9990 | 0.866296 | 0.200000 |
| 3 | RepeatedKFold Cross Validation | Tree Model with gini criteria | 0.9989 | 0.833033 | 1.000000 |
| 4 | RepeatedKFold Cross Validation | Tree Model with entropy criteria | 0.9992 | 0.899750 | 1.000000 |
| 5 | RepeatedKFold Cross Validation | Random Forest | 0.9997 | 0.965381 | 0.300000 |
| 6 | RepeatedKFold Cross Validation | XGBoost | 0.9996 | 0.936365 | 0.283122 |
| 7 | RepeatedKFold Cross Validation | SVM | 0.9981 | 0.724487 | 0.003482 |
| 8 | RepeatedKFold Cross Validation | ANN | 0.9984 | 0.927278 | 0.000521 |
| 9 | StratifiedKFold Cross Validation | Logistic Regression with L2 Regularisation | 0.9992 | 0.956638 | 0.002113 |
| 10 | StratifiedKFold Cross Validation | Logistic Regression with L1 Regularisation | 0.9991 | 0.892559 | 0.301463 |
| 11 | StratifiedKFold Cross Validation | KNN | 0.9990 | 0.823250 | 0.200000 |
| 12 | StratifiedKFold Cross Validation | Tree Model with gini criteria | 0.9985 | 0.823079 | 1.000000 |
| 13 | StratifiedKFold Cross Validation | Tree Model with entropy criteria | 0.9990 | 0.823329 | 1.000000 |
| 14 | StratifiedKFold Cross Validation | Random Forest | 0.9995 | 0.937102 | 0.010000 |
| 15 | StratifiedKFold Cross Validation | XGBoost | 0.9993 | 0.971463 | 0.001395 |
| 16 | StratifiedKFold Cross Validation | SVM | 0.9975 | 0.634779 | 0.003424 |
| 17 | StratifiedKFold Cross Validation | ANN | 0.9985 | 0.925528 | 0.000077 |

Results for cross validation with StratifiedKFold:

- Looking at the ROC value we have Logistic Regression with L2 Regularisation has provided best results for cross validation with StratifiedKFold technique

Conclusion:

- As the results show Logistic Regression with L2 Regularisation for StratifiedKFold cross validation provided best results

Proceed with the model which shows the best result

- Apply the best hyperparameter on the model
- Predict on the test dataset

```
In [378]: # Logistic Regression
from sklearn import linear_model #import the package
from sklearn.model_selection import KFold

num_C = list(np.power(10.0, np.arange(-10, 10)))
cv_num = KFold(n_splits=10, shuffle=True, random_state=42)

clf = linear_model.LogisticRegressionCV(
    Cs= num_C
    ,penalty='l2'
    ,scoring='roc_auc'
    ,cv=cv_num
    ,random_state=42
    ,max_iter=10000
    ,fit_intercept=True
    ,solver='newton-cg'
    ,tol=10
)

clf.fit(X_train_SKF_cv, y_train_SKF_cv)
print ('Max auc_roc for l2:', clf.scores_[1].mean(axis=0).max())

print("Parameters for l2 regularisations")
print(clf.coef_)
print(clf.intercept_)
print(clf.scores_)

#find predicted values
y_pred_l2 = clf.predict(X_test)

#Find predicted probabilities
y_pred_probs_l2 = clf.predict_proba(X_test)[:,:1]

# Accuracy of L2/L1 models
Accuracy_l2 = metrics.accuracy_score(y_pred=y_pred_l2, y_true=y_test)

print("Accuracy of Logistic model with l2 regularisation : {}".format(Accuracy_l2))
```

```
Max auc_roc for l2: 0.9567308642474954
Parameters for l2 regularisations
[[-1.31491217e-02  3.07463695e-02 -1.23758286e-01  2.90976907e-01
  1.00763702e-01 -5.25175254e-02  4.50978185e-02 -1.31263201e-01
 -1.08854534e-01 -1.56738327e-01  1.83509896e-01 -2.25403516e-01
 -5.45158708e-02 -4.16368675e-01 -5.45595030e-02 -7.19437239e-02
 -1.10316217e-01 -2.06331993e-02  2.04222421e-02 -2.72144122e-02
 4.33597690e-02  6.29205995e-02  3.69527095e-02 -1.48824898e-02
 -1.39901365e-02  2.51950981e-03 -4.00624298e-03 -7.95897195e-02
  3.42612078e-04  1.20704152e-02]]
[-7.64242687]
{1: array([[0.65944862, 0.65944862, 0.66107769, 0.67919799, 0.77882206,
  0.84032581, 0.88531328, 0.92263158, 0.9356391 , 0.93924812,
  0.93924812, 0.93924812, 0.93924812, 0.93924812, 0.93924812],
 [0.55307962, 0.55307962, 0.55474879, 0.58771491, 0.73147221,
  0.84848105, 0.87819229, 0.90110165, 0.91103322, 0.89542647,
  0.89542647, 0.89542647, 0.89542647, 0.89542647, 0.89542647],
 [0.49772817, 0.49772817, 0.49898036, 0.52584881, 0.67553934,
  0.87138206, 0.94204143, 0.97273801, 0.98007227, 0.98443705,
  0.98443705, 0.98443705, 0.98443705, 0.98443705, 0.98443705],
 [0.58863726, 0.58860148, 0.59024722, 0.61833208, 0.79202891,
  0.9631498 , 0.9797145 , 0.92626382, 0.90862581, 0.91062932,
  0.91062932, 0.91062932, 0.91062932, 0.91062932, 0.91062932],
 [0.89706987, 0.89692676, 0.89774963, 0.90143465, 0.92533362,
  0.9622196 , 0.93864262, 0.95302494, 0.98339952, 0.98164645,
  0.98164645, 0.98164645, 0.98164645, 0.98164645, 0.98164645],
 [0.50973823, 0.50973823, 0.51111598, 0.53710546, 0.66846192,
  0.79950526, 0.90812876, 0.97263277, 0.96940757, 0.96157941,
  0.9332728 , 0.9332728 , 0.9332728 , 0.9332728 , 0.9332728 ],
 [0.51582833, 0.51582833, 0.51607858, 0.52152152, 0.57645145,
  0.76295045, 0.90546797, 0.90934685, 0.92880038 , 0.87825325,
  0.87825325, 0.87825325, 0.87825325, 0.87825325, 0.87825325],
 [0.41166166, 0.41166166, 0.41278779, 0.44575826, 0.6265015 ,
  0.80974725, 0.90040004 , 0.94513263, 0.95952202, 0.95658158,
  0.95658158, 0.94331832, 0.94331832, 0.94331832, 0.94331832],
 [0.544041 , 0.544041 , 0.55550785, 0.57096347, 0.69085185,
  0.88064828, 0.9584988 , 0.98533147, 0.99080534, 0.98966048,
  0.98966048, 0.98966048, 0.98966048, 0.98966048, 0.98966048],
 [0.33032807, 0.33022074, 0.33458552, 0.38728489, 0.68140675,
  0.94451004, 0.99745984, 0.99992845, 1. , 1. ,
  1. , 1. , 1. , 1. , 1. ]])
Accuracy of Logistic model with l2 regularisation : 0.9995
12 roc_value: 0.9785411450509097
12 threshold: 0.008174170065796222
```

```
In [379]: # Checking for the coefficient values
clf.coef_
```

```
Out[379]: array([[ -1.31491217e-02,   3.07463695e-02,  -1.23758286e-01,
    2.90976907e-01,   1.00763702e-01,  -5.25175254e-02,
    4.50978185e-02,  -1.31263201e-01,  -1.08854534e-01,
   -1.56738327e-01,   1.83509896e-01,  -2.25403516e-01,
   -5.45158708e-02,  -4.16368675e-01,  -5.45595030e-02,
   -7.19437239e-02,  -1.10316217e-01,  -2.06331993e-02,
    2.04222421e-02,  -2.72144122e-02,   4.33597690e-02,
    6.29205995e-02,   3.69527095e-02,  -1.48824898e-02,
   -1.39901365e-02,   2.51950981e-03,  -4.00624298e-03,
   -7.95897195e-02,   3.42612078e-04,   1.20704152e-02]])
```

```
In [380]: # Creating a dataframe with the coefficient values
coefficients = pd.concat([pd.DataFrame(X.columns),pd.DataFrame(np.transpose(clf.coef_))], axis = 1)
coefficients.columns = ['Feature', 'Importance Coefficient']
```

```
In [381]: coefficients
```


```
Out[381]:
```

| | Feature | Importance Coefficient |
|----|-----------|------------------------|
| 0 | V1 | -0.013149 |
| 1 | V2 | 0.030746 |
| 2 | V3 | -0.123758 |
| 3 | V4 | 0.290977 |
| 4 | V5 | 0.100764 |
| 5 | V6 | -0.052518 |
| 6 | V7 | 0.045098 |
| 7 | V8 | -0.131263 |
| 8 | V9 | -0.108855 |
| 9 | V10 | -0.156738 |
| 10 | V11 | 0.183510 |
| 11 | V12 | -0.225404 |
| 12 | V13 | -0.054516 |
| 13 | V14 | -0.416369 |
| 14 | V15 | -0.054560 |
| 15 | V16 | -0.071944 |
| 16 | V17 | -0.110316 |
| 17 | V18 | -0.020633 |
| 18 | V19 | 0.020422 |
| 19 | V20 | -0.027214 |
| 20 | V21 | 0.043360 |
| 21 | V22 | 0.062921 |
| 22 | V23 | 0.036953 |
| 23 | V24 | -0.014882 |
| 24 | V25 | -0.013990 |
| 25 | V26 | 0.002520 |
| 26 | V27 | -0.004006 |
| 27 | V28 | -0.079590 |
| 28 | Amount | 0.000343 |
| 29 | Time_Hour | 0.012070 |

Print the important features of the best model to understand the dataset

- This will not give much explanation on the already transformed dataset
- But it will help us in understanding if the dataset is not PCA transformed

```
In [382]: # Plotting the coefficient values
plt.figure(figsize=(20,5))
sns.barplot(x='Feature', y='Importance Coefficient', data=coefficients)
plt.title("Logistic Regression with L2 Regularisation Feature Importance", fontsize=18)
plt.show()
```



- Hence it implies that V4, v5,V11 has + ve importance whereas V10, V12, V14 seems to have -ve impact on the predictions

Model building with balancing Classes

Perform class balancing with:

- Random Oversampling
- SMOTE
- ADASYN

Oversampling with RandomOverSampler with StratifiedKFold Cross Validation

- **Random Oversampling:**
 - Random Oversampling is a technique for handling imbalanced datasets by duplicating instances from the minority class to balance the class distribution.
 - This method increases the number of samples in the minority class to compensate for the class imbalance, but it can also lead to overfitting, as it increases the likelihood of duplicating instances in the training set.



```
In [383]: # Creating the dataset with RandomOverSampler and StratifiedKFold
from sklearn.model_selection import StratifiedKFold
from imblearn.over_sampling import RandomOverSampler

skf = StratifiedKFold(n_splits=5, random_state=None)

for fold, (train_index, test_index) in enumerate(skf.split(X,y), 1):
    # print(fold, train_index, test_index)
    X_train = X.iloc[train_index]
    y_train = y.iloc[train_index]
    X_test = X.iloc[test_index]
    y_test = y.iloc[test_index]
    ROS = RandomOverSampler(sampling_strategy=0.5)
    X_over, y_over= ROS.fit_resample(X_train, y_train)

X_over = pd.DataFrame(data=X_over, columns=cols)
```

```

In [384]: Data_Imbalance_Handling = "Random Oversampling with StratifiedKFold CV "
#Run Logistic Regression with L1 And L2 Regularisation
print("Logistic Regression with L1 And L2 Regularisation")
start_time = time.time()
df_Results = buildAndRunLogisticModels(df_Results , Data_Imbalance_Handling , X_over, y_over, X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*60 )

#Run KNN Model
print("KNN Model")
start_time = time.time()
df_Results = buildAndRunKNNModels(df_Results , Data_Imbalance_Handling,X_over, y_over, X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*60 )

#Run Decision Tree Models with 'gini' & 'entropy' criteria
print("Decision Tree Models with 'gini' & 'entropy' criteria")
start_time = time.time()
df_Results = buildAndRunTreeModels(df_Results , Data_Imbalance_Handling,X_over, y_over, X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*60 )


#Run Random Forest Model
print("Random Forest Model")
start_time = time.time()
df_Results = buildAndRunRandomForestModels(df_Results , Data_Imbalance_Handling,X_over, y_over, X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*60 )

#Run XGBoost Model
print("XGBoost Model")
start_time = time.time()
df_Results = buildAndRunXGBoostModels(df_Results , Data_Imbalance_Handling,X_over, y_over, X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*60 )

#Run ANN Model
print("ANN Model")
start_time = time.time()
df_Results = buildAndgetmetricsANN(df_Results, Data_Imbalance_Handling, X_train_Smote, y_train_Smote , X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*80 )


```

Accuracy of Logistic model with l2 regularisation : 0.9808
Confusion Matrix



| classification Report | | recall | f1-score | support |
|-----------------------|-----------|--------|----------|---------|
| | precision | | | |
| 0 | 1.00 | 0.98 | 0.99 | 9983 |
| 1 | 0.07 | 0.82 | 0.13 | 17 |
| accuracy | | | 0.98 | 10000 |
| macro avg | 0.53 | 0.90 | 0.56 | 10000 |
| weighted avg | 1.00 | 0.98 | 0.99 | 10000 |


Accuarcy of Logistic model with l1 regularisation : 0.9983
Confusion Matrix




Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.00 | 0.00 | 0.00 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.50 | 0.50 | 0.50 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

12 roc_value: 0.9630784097671925
 12 threshold: 0.3837371155321852
 ROC for the test dataset 96.3%




11 roc_value: 0.5
 11 threshold: 1.5
 ROC for the test dataset 50.0%



Time Taken by Model: --- 99.70262217521667 seconds ---


KNN Model
 model score
 0.9986
 Confusion Matrix



Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.58 | 0.65 | 0.61 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.79 | 0.82 | 0.81 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |


KNN roc_value: 0.823164084826558
 KNN threshold: 0.6
 ROC for the test dataset 82.3%



Time Taken by Model: --- 76.98786473274231 seconds ---

Decision Tree Models with 'gini' & 'entropy' criteria
gini score: 0.9987

Confusion Matrix




| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.62 | 0.59 | 0.61 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.81 | 0.79 | 0.80 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

gini tree_roc_value: 0.7938171361903471


Tree threshold: 1.0

ROC for the test dataset 79.4%



entropy score: 0.9991

Confusion Matrix




| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.79 | 0.65 | 0.71 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.89 | 0.82 | 0.85 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

entropy tree_roc_value: 0.8233791563304677

Tree threshold: 1.0

ROC for the test dataset 82.3%




Time Taken by Model: --- 2.96124267578125 seconds ---

Random Forest Model

Model Accuracy: 0.9994

Confusion Matrix




| Classification Report | | | | |
|-----------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.92 | 0.71 | 0.80 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.96 | 0.85 | 0.90 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

Random Forest roc_value: 0.9378914743298901

Random Forest threshold: 0.02

ROC for the test dataset 93.8%




Time Taken by Model: --- 17.280887126922607 seconds ---

XGBoost Model

Model Accuracy: 0.9992

Confusion Matrix




| Classification Report | | | | |
|-----------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.76 | 0.76 | 0.76 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.88 | 0.88 | 0.88 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

XGboost roc_value: 0.9714750369746215


XGBoost threshold: 0.13544505834579468

ROC for the test dataset 97.1%



Time Taken by Model: --- 11.824404954910278 seconds ---


 ANN Model
 Epoch 1/5
 2130/2130 [=====] - 7s 3ms/step - loss: 0.1501 - accuracy: 0.9519
 Epoch 2/5
 2130/2130 [=====] - 6s 3ms/step - loss: 0.0373 - accuracy: 0.9907
 Epoch 3/5
 2130/2130 [=====] - 5s 3ms/step - loss: 0.0187 - accuracy: 0.9945
 Epoch 4/5
 2130/2130 [=====] - 9s 4ms/step - loss: 0.0177 - accuracy: 0.9959
 Epoch 5/5
 2130/2130 [=====] - 5s 3ms/step - loss: 0.0199 - accuracy: 0.9953
 999/999 [=====] - 4s 4ms/step - loss: 0.0079 - accuracy: 0.9978
 313/313 [=====] - 1s 2ms/step
 accuracy_score : 0.9991
 Confusion Matrix



classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.72 | 0.76 | 0.74 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.86 | 0.88 | 0.87 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

ANN roc_value: 0.9469333160490481
 ANN threshold: 7.360094514297089e-06
 ROC for the test dataset 94.7%



Time Taken by Model: --- 38.464799880981445 seconds ---

```
In [385]: # Checking the df_result dataframe which contains consolidated results of all the runs
df_Results
```

```
Out[385]:
```

| | Methodology | Model | Accuracy | roc_value | threshold |
|----|---|--|----------|-----------|-----------|
| 0 | RepeatedKFold Cross Validation | Logistic Regression with L2 Regularisation | 0.9990 | 0.950886 | 0.030690 |
| 1 | RepeatedKFold Cross Validation | Logistic Regression with L1 Regularisation | 0.9991 | 0.933580 | 0.057833 |
| 2 | RepeatedKFold Cross Validation | KNN | 0.9990 | 0.866296 | 0.200000 |
| 3 | RepeatedKFold Cross Validation | Tree Model with gini criteria | 0.9989 | 0.833033 | 1.000000 |
| 4 | RepeatedKFold Cross Validation | Tree Model with entropy criteria | 0.9992 | 0.899750 | 1.000000 |
| 5 | RepeatedKFold Cross Validation | Random Forest | 0.9997 | 0.965381 | 0.300000 |
| 6 | RepeatedKFold Cross Validation | XGBoost | 0.9996 | 0.936365 | 0.283122 |
| 7 | RepeatedKFold Cross Validation | SVM | 0.9981 | 0.724487 | 0.003482 |
| 8 | RepeatedKFold Cross Validation | ANN | 0.9984 | 0.927278 | 0.000521 |
| 9 | StratifiedKFold Cross Validation | Logistic Regression with L2 Regularisation | 0.9992 | 0.956638 | 0.002113 |
| 10 | StratifiedKFold Cross Validation | Logistic Regression with L1 Regularisation | 0.9991 | 0.892559 | 0.301463 |
| 11 | StratifiedKFold Cross Validation | KNN | 0.9990 | 0.823250 | 0.200000 |
| 12 | StratifiedKFold Cross Validation | Tree Model with gini criteria | 0.9985 | 0.823079 | 1.000000 |
| 13 | StratifiedKFold Cross Validation | Tree Model with entropy criteria | 0.9990 | 0.823329 | 1.000000 |
| 14 | StratifiedKFold Cross Validation | Random Forest | 0.9995 | 0.937102 | 0.010000 |
| 15 | StratifiedKFold Cross Validation | XGBoost | 0.9993 | 0.971463 | 0.001395 |
| 16 | StratifiedKFold Cross Validation | SVM | 0.9975 | 0.634779 | 0.003424 |
| 17 | StratifiedKFold Cross Validation | ANN | 0.9985 | 0.925528 | 0.000077 |
| 18 | Random Oversampling with StratifiedKFold CV | Logistic Regression with L2 Regularisation | 0.9808 | 0.963078 | 0.383737 |
| 19 | Random Oversampling with StratifiedKFold CV | Logistic Regression with L1 Regularisation | 0.9983 | 0.500000 | 1.500000 |
| 20 | Random Oversampling with StratifiedKFold CV | KNN | 0.9986 | 0.823164 | 0.600000 |
| 21 | Random Oversampling with StratifiedKFold CV | Tree Model with gini criteria | 0.9987 | 0.793817 | 1.000000 |
| 22 | Random Oversampling with StratifiedKFold CV | Tree Model with entropy criteria | 0.9991 | 0.823379 | 1.000000 |
| 23 | Random Oversampling with StratifiedKFold CV | Random Forest | 0.9994 | 0.937891 | 0.020000 |
| 24 | Random Oversampling with StratifiedKFold CV | XGBoost | 0.9992 | 0.971475 | 0.135445 |
| 25 | Random Oversampling with StratifiedKFold CV | ANN | 0.9991 | 0.946933 | 0.000007 |


Results for Random Oversampling with StratifiedKFold technique:

- Looking at the Accuracy and ROC value we have XGBoost which has provided best results for Random Oversampling and StratifiedKFold technique

Oversampling with SMOTE Oversampling

SMOTE (Synthetic Minority Over-sampling Technique):

- SMOTE (Synthetic Minority Over-sampling Technique) is a data augmentation technique for handling imbalanced datasets by creating synthetic samples for the minority class.
- SMOTE generates new instances for the minority class by interpolating between existing instances and their nearest neighbors, creating a more diverse representation of the minority class. This can help balance the class distribution and prevent overfitting due to random oversampling.



```
In [386]: # Creating dataframe with Smote and StratifiedKFold
from sklearn.model_selection import StratifiedKFold
from imblearn import over_sampling

skf = StratifiedKFold(n_splits=5, random_state=None)

for fold, (train_index, test_index) in enumerate(skf.split(X,y), 1):
    X_train = X.iloc[train_index]
    y_train = y.iloc[train_index]
    X_test = X.iloc[test_index]
    y_test = y.iloc[test_index]
    SMOTE = over_sampling.SMOTE(random_state=0)
    X_train_Smote, y_train_Smote= SMOTE.fit_resample(X_train, y_train)

X_train_Smote = pd.DataFrame(data=X_train_Smote, columns=cols)
```



```

In [387]: Data_Imbalance_Handling = "SMOTE Oversampling with StratifiedKFold CV "
#Run Logistic Regression with L1 And L2 Regularisation
print("Logistic Regression with L1 And L2 Regularisation")
start_time = time.time()
df_Results = buildAndRunLogisticModels(df_Results, Data_Imbalance_Handling, X_train_Smote, y_train_Smote , X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*80 )

#Run KNN Model
print("KNN Model")
start_time = time.time()
df_Results = buildAndRunKNNModels(df_Results, Data_Imbalance_Handling, X_train_Smote, y_train_Smote , X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*80 )

#Run Decision Tree Models with 'gini' & 'entropy' criteria
print("Decision Tree Models with 'gini' & 'entropy' criteria")
start_time = time.time()
df_Results = buildAndRunTreeModels(df_Results, Data_Imbalance_Handling, X_train_Smote, y_train_Smote , X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*80 )


#Run Random Forest Model
print("Random Forest Model")
start_time = time.time()
df_Results = buildAndRunRandomForestModels(df_Results, Data_Imbalance_Handling, X_train_Smote, y_train_Smote , X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*80 )

#Run XGBoost Model
print("XGBoost Model")
start_time = time.time()
df_Results = buildAndRunXGBoostModels(df_Results, Data_Imbalance_Handling, X_train_Smote, y_train_Smote , X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*80 )

#Run ANN Model
print("ANN Model")
start_time = time.time()
df_Results = buildAndgetmetricsANN(df_Results, Data_Imbalance_Handling, X_train_Smote, y_train_Smote , X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*80 )


```

Accuracy of Logistic model with l2 regularisation : 0.97
Confusion Matrix



| classification Report | | recall | f1-score | support |
|-----------------------|-----------|--------|----------|---------|
| | precision | | | |
| 0 | 1.00 | 0.97 | 0.98 | 9983 |
| 1 | 0.05 | 0.88 | 0.09 | 17 |
| accuracy | | | 0.97 | 10000 |
| macro avg | 0.52 | 0.93 | 0.54 | 10000 |
| weighted avg | 1.00 | 0.97 | 0.98 | 10000 |


Accuracy of Logistic model with l1 regularisation : 0.9983
Confusion Matrix




Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.00 | 0.00 | 0.00 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.50 | 0.50 | 0.50 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

12 roc_value: 0.9636470234693096
 12 threshold: 0.3707247623650453
 ROC for the test dataset 96.4%




11 roc_value: 0.5
 11 threshold: 1.5
 ROC for the test dataset 50.0%



Time Taken by Model: --- 118.28125309944153 seconds ---


KNN Model
 model score
 0.9925
 Confusion Matrix



Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.99 | 1.00 | 9983 |
| 1 | 0.15 | 0.71 | 0.24 | 17 |
| accuracy | | | 0.99 | 10000 |
| macro avg | 0.57 | 0.85 | 0.62 | 10000 |
| weighted avg | 1.00 | 0.99 | 0.99 | 10000 |

KNN roc_value: 0.8501540854747189
 KNN threshold: 1.0
 ROC for the test dataset 85.0%




Time Taken by Model: --- 106.98998665809631 seconds ---

Decision Tree Models with 'gini' & 'entropy' criteria

gini score: 0.9955

Confusion Matrix




| Classification Report | | | | |
|-----------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.17 | 0.41 | 0.24 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.58 | 0.70 | 0.62 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

gini tree_roc_value: 0.704129372875064


Tree threshold: 1.0

ROC for the test dataset 70.4%



entropy score: 0.9974

Confusion Matrix




| Classification Report | | | | |
|-----------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.35 | 0.65 | 0.46 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.68 | 0.82 | 0.73 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

entropy tree_roc_value: 0.8225277088697844

Tree threshold: 1.0

ROC for the test dataset 82.3%




Time Taken by Model: --- 9.980096578598022 seconds ---

Random Forest Model

Model Accuracy: 0.9995

Confusion Matrix




Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.93 | 0.76 | 0.84 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.96 | 0.88 | 0.92 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

Random Forest roc_value: 0.9589478584181342

Random Forest threshold: 0.18

ROC for the test dataset 95.9%




Time Taken by Model: --- 43.73682975769043 seconds ---

XGBoost Model

Model Accuracy: 0.9948

Confusion Matrix




Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.99 | 1.00 | 9983 |
| 1 | 0.23 | 0.88 | 0.37 | 17 |
| accuracy | | | 0.99 | 10000 |
| macro avg | 0.62 | 0.94 | 0.68 | 10000 |
| weighted avg | 1.00 | 0.99 | 1.00 | 10000 |

XGboost roc_value: 0.9464030027517367


XGBoost threshold: 0.5603106617927551

ROC for the test dataset 94.6%




Time Taken by Model: --- 21.104939460754395 seconds ---

 ANN Model
 Epoch 1/5
 5325/5325 [=====] - 15s 3ms/step - loss: 0.2444 - accuracy: 0.9168
 Epoch 2/5
 5325/5325 [=====] - 15s 3ms/step - loss: 0.0891 - accuracy: 0.9701
 Epoch 3/5
 5325/5325 [=====] - 15s 3ms/step - loss: 0.0657 - accuracy: 0.9789
 Epoch 4/5
 5325/5325 [=====] - 17s 3ms/step - loss: 0.0516 - accuracy: 0.9838
 Epoch 5/5
 5325/5325 [=====] - 14s 3ms/step - loss: 0.0393 - accuracy: 0.9877
 2496/2496 [=====] - 8s 3ms/step - loss: 0.0293 - accuracy: 0.9910
 313/313 [=====] - 1s 1ms/step
 accuracy_score : 0.9941
 Confusion Matrix



| classification Report | | | | |
|-----------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 1.00 | 0.99 | 1.00 | 9983 |
| 1 | 0.20 | 0.82 | 0.32 | 17 |
| accuracy | | | 0.99 | 10000 |
| macro avg | 0.60 | 0.91 | 0.66 | 10000 |
| weighted avg | 1.00 | 0.99 | 1.00 | 10000 |

ANN roc_value: 0.944962318294041
 ANN threshold: 0.06009089574217796
 ROC for the test dataset 94.5%



Time Taken by Model: --- 87.03782868385315 seconds ---

```
In [388]: # Checking the df_result dataframe which contains consolidated results of all the runs
df_Results
```

```
Out[388]:
```

| | Methodology | Model | Accuracy | roc_value | threshold |
|----|---|--|----------|-----------|-----------|
| 0 | RepeatedKFold Cross Validation | Logistic Regression with L2 Regularisation | 0.9990 | 0.950886 | 0.030690 |
| 1 | RepeatedKFold Cross Validation | Logistic Regression with L1 Regularisation | 0.9991 | 0.933580 | 0.057833 |
| 2 | RepeatedKFold Cross Validation | KNN | 0.9990 | 0.866296 | 0.200000 |
| 3 | RepeatedKFold Cross Validation | Tree Model with gini criteria | 0.9989 | 0.833033 | 1.000000 |
| 4 | RepeatedKFold Cross Validation | Tree Model with entropy criteria | 0.9992 | 0.899750 | 1.000000 |
| 5 | RepeatedKFold Cross Validation | Random Forest | 0.9997 | 0.965381 | 0.300000 |
| 6 | RepeatedKFold Cross Validation | XGBoost | 0.9996 | 0.936365 | 0.283122 |
| 7 | RepeatedKFold Cross Validation | SVM | 0.9981 | 0.724487 | 0.003482 |
| 8 | RepeatedKFold Cross Validation | ANN | 0.9984 | 0.927278 | 0.000521 |
| 9 | StratifiedKFold Cross Validation | Logistic Regression with L2 Regularisation | 0.9992 | 0.956638 | 0.002113 |
| 10 | StratifiedKFold Cross Validation | Logistic Regression with L1 Regularisation | 0.9991 | 0.892559 | 0.301463 |
| 11 | StratifiedKFold Cross Validation | KNN | 0.9990 | 0.823250 | 0.200000 |
| 12 | StratifiedKFold Cross Validation | Tree Model with gini criteria | 0.9985 | 0.823079 | 1.000000 |
| 13 | StratifiedKFold Cross Validation | Tree Model with entropy criteria | 0.9990 | 0.823329 | 1.000000 |
| 14 | StratifiedKFold Cross Validation | Random Forest | 0.9995 | 0.937102 | 0.010000 |
| 15 | StratifiedKFold Cross Validation | XGBoost | 0.9993 | 0.971463 | 0.001395 |
| 16 | StratifiedKFold Cross Validation | SVM | 0.9975 | 0.634779 | 0.003424 |
| 17 | StratifiedKFold Cross Validation | ANN | 0.9985 | 0.925528 | 0.000077 |
| 18 | Random Oversampling with StratifiedKFold CV | Logistic Regression with L2 Regularisation | 0.9808 | 0.963078 | 0.383737 |
| 19 | Random Oversampling with StratifiedKFold CV | Logistic Regression with L1 Regularisation | 0.9983 | 0.500000 | 1.500000 |
| 20 | Random Oversampling with StratifiedKFold CV | KNN | 0.9986 | 0.823164 | 0.600000 |
| 21 | Random Oversampling with StratifiedKFold CV | Tree Model with gini criteria | 0.9987 | 0.793817 | 1.000000 |
| 22 | Random Oversampling with StratifiedKFold CV | Tree Model with entropy criteria | 0.9991 | 0.823379 | 1.000000 |
| 23 | Random Oversampling with StratifiedKFold CV | Random Forest | 0.9994 | 0.937891 | 0.020000 |
| 24 | Random Oversampling with StratifiedKFold CV | XGBoost | 0.9992 | 0.971475 | 0.135445 |
| 25 | Random Oversampling with StratifiedKFold CV | ANN | 0.9991 | 0.946933 | 0.000007 |
| 26 | SMOTE Oversampling with StratifiedKFold CV | Logistic Regression with L2 Regularisation | 0.9700 | 0.963647 | 0.370725 |
| 27 | SMOTE Oversampling with StratifiedKFold CV | Logistic Regression with L1 Regularisation | 0.9983 | 0.500000 | 1.500000 |
| 28 | SMOTE Oversampling with StratifiedKFold CV | KNN | 0.9925 | 0.850154 | 1.000000 |
| 29 | SMOTE Oversampling with StratifiedKFold CV | Tree Model with gini criteria | 0.9955 | 0.704129 | 1.000000 |
| 30 | SMOTE Oversampling with StratifiedKFold CV | Tree Model with entropy criteria | 0.9974 | 0.822528 | 1.000000 |
| 31 | SMOTE Oversampling with StratifiedKFold CV | Random Forest | 0.9995 | 0.958948 | 0.180000 |
| 32 | SMOTE Oversampling with StratifiedKFold CV | XGBoost | 0.9948 | 0.946403 | 0.560311 |
| 33 | SMOTE Oversampling with StratifiedKFold CV | ANN | 0.9941 | 0.944962 | 0.060091 |


Results for SMOTE Oversampling with StratifiedKFold:

- Looking at Accuracy and ROC value we have XGBoost which has provided best results for SMOTE Oversampling with StratifiedKFold technique

Oversampling with ADASYN Oversampling

ADASYN (Adaptive Synthetic Sampling):

- ADASYN (Adaptive Synthetic Sampling) is a data augmentation technique for handling imbalanced datasets by creating synthetic samples for the minority class with adaptive weighting.
- Unlike SMOTE, which generates synthetic samples for all minority class instances with equal weight, ADASYN assigns higher weights to minority class instances that are harder to classify, leading to a more balanced distribution of synthetic samples. This can help improve the performance of classification models on imbalanced datasets.



```
In [389]: # Creating dataframe with ADASYN and StratifiedKFold
from sklearn.model_selection import StratifiedKFold
from imblearn import over_sampling

skf = StratifiedKFold(n_splits=5, random_state=None)

for fold, (train_index, test_index) in enumerate(skf.split(X,y), 1):
    X_train = X.iloc[train_index]
    y_train = y.iloc[train_index]
    X_test = X.iloc[test_index]
    y_test = y.iloc[test_index]
    ADASYN = over_sampling.ADasYN(random_state=0)
    X_train_ADASYN, y_train_ADASYN= ADASYN.fit_resample(X_train, y_train)

X_train_ADASYN = pd.DataFrame(data=X_train_ADASYN, columns=cols)
```



```
In [390]: Data_Imbalance_Handling = "ADASYN Oversampling with StratifiedKFold CV"
#Run Logistic Regression with L1 And L2 Regularisation
print("Logistic Regression with L1 And L2 Regularisation")
start_time = time.time()
df_Results = buildAndRunLogisticModels(df_Results, Data_Imbalance_Handling, X_train_ADASYN, y_train_ADASYN, X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*80)

#Run KNN Model
print("KNN Model")
start_time = time.time()
df_Results = buildAndRunKNNModels(df_Results, Data_Imbalance_Handling, X_train_ADASYN, y_train_ADASYN, X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*80)


#Run Decision Tree Models with 'gini' & 'entropy' criteria
print("Decision Tree Models with 'gini' & 'entropy' criteria")
start_time = time.time()
df_Results = buildAndRunTreeModels(df_Results, Data_Imbalance_Handling, X_train_ADASYN, y_train_ADASYN, X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*80)

#Run Random Forest Model
print("Random Forest Model")
start_time = time.time()
df_Results = buildAndRunRandomForestModels(df_Results, Data_Imbalance_Handling, X_train_ADASYN, y_train_ADASYN, X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*80)

#Run XGBoost Model
print("XGBoost Model")
start_time = time.time()
df_Results = buildAndRunXGBoostModels(df_Results, Data_Imbalance_Handling, X_train_ADASYN, y_train_ADASYN, X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*80)


#Run ANN Model
print("ANN Model")
start_time = time.time()
df_Results = buildAndgetmetricsANN(df_Results, Data_Imbalance_Handling, X_train_Smote, y_train_Smote, X_test, y_test)
print("Time Taken by Model: --- %s seconds ---" % (time.time() - start_time))
print('-'*80)
```

Accuracy of Logistic model with l2 regularisation : 0.9567



| classification Report | | recall | f1-score | support |
|-----------------------|-----------|--------|----------|---------|
| | precision | | | |
| 0 | 1.00 | 0.96 | 0.98 | 9983 |
| 1 | 0.04 | 0.94 | 0.07 | 17 |
| accuracy | | | 0.96 | 10000 |
| macro avg | 0.52 | 0.95 | 0.52 | 10000 |
| weighted avg | 1.00 | 0.96 | 0.98 | 10000 |


Accuracy of Logistic model with l1 regularisation : 0.9983
Confusion Matrix




Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.00 | 0.00 | 0.00 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.50 | 0.50 | 0.50 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

12 roc_value: 0.964887367348021
 12 threshold: 0.535068585528712
 ROC for the test dataset 96.5%




11 roc_value: 0.5
 11 threshold: 1.5
 ROC for the test dataset 50.0%



Time Taken by Model: --- 116.42573523521423 seconds ---


KNN Model
 model score
 0.9919
 Confusion Matrix



Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.99 | 1.00 | 9983 |
| 1 | 0.14 | 0.71 | 0.23 | 17 |
| accuracy | | | 0.99 | 10000 |
| macro avg | 0.57 | 0.85 | 0.61 | 10000 |
| weighted avg | 1.00 | 0.99 | 0.99 | 10000 |

KNN roc_value: 0.8500804308500923
 KNN threshold: 1.0
 ROC for the test dataset 85.0%




Time Taken by Model: --- 90.96645188331604 seconds ---

Decision Tree Models with 'gini' & 'entropy' criteria

gini score: 0.9971

Confusion Matrix




Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.27 | 0.41 | 0.33 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.63 | 0.70 | 0.66 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

gini tree_roc_value: 0.7049307351910011


Tree threshold: 1.0

ROC for the test dataset 70.5%



entropy score: 0.997

Confusion Matrix




Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.30 | 0.59 | 0.40 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.65 | 0.79 | 0.70 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

entropy tree_roc_value: 0.792965688729664

Tree threshold: 1.0

ROC for the test dataset 79.3%




Time Taken by Model: --- 8.870583772659302 seconds ---

Random Forest Model

Model Accuracy: 0.9993


Confusion Matrix



| Classification Report | | | | |
|-----------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.86 | 0.71 | 0.77 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.93 | 0.85 | 0.89 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

Random Forest roc_value: 0.9598876914283694

Random Forest threshold: 0.15
ROC for the test dataset 96.0%




Time Taken by Model: --- 44.52493453025818 seconds ---

XGBoost Model

Model Accuracy: 0.9941


Confusion Matrix



| Classification Report | | | | |
|-----------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 1.00 | 0.99 | 1.00 | 9983 |
| 1 | 0.21 | 0.88 | 0.34 | 17 |
| accuracy | | | 0.99 | 10000 |
| macro avg | 0.60 | 0.94 | 0.67 | 10000 |
| weighted avg | 1.00 | 0.99 | 1.00 | 10000 |


XGboost roc_value: 0.9527667623194724

XGBoost threshold: 0.575517475605011
ROC for the test dataset 95.3%



Time Taken by Model: --- 20.68782329559326 seconds ---


 ANN Model
 Epoch 1/5
 5325/5325 [=====] - 16s 3ms/step - loss: 0.1769 - accuracy: 0.9382
 Epoch 2/5
 5325/5325 [=====] - 15s 3ms/step - loss: 0.0594 - accuracy: 0.9793
 Epoch 3/5
 5325/5325 [=====] - 15s 3ms/step - loss: 0.0404 - accuracy: 0.9861
 Epoch 4/5
 5325/5325 [=====] - 16s 3ms/step - loss: 0.0291 - accuracy: 0.9909
 Epoch 5/5
 5325/5325 [=====] - 15s 3ms/step - loss: 0.0249 - accuracy: 0.9928
 2496/2496 [=====] - 8s 3ms/step - loss: 0.0161 - accuracy: 0.9956
 313/313 [=====] - 1s 1ms/step
 accuracy_score : 0.998
 Confusion Matrix



Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 9983 |
| 1 | 0.45 | 0.76 | 0.57 | 17 |
| accuracy | | | 1.00 | 10000 |
| macro avg | 0.72 | 0.88 | 0.78 | 10000 |
| weighted avg | 1.00 | 1.00 | 1.00 | 10000 |

ANN roc_value: 0.9535327704155887
 ANN threshold: 0.0017699062591418624
 ROC for the test dataset 95.4%



Time Taken by Model: --- 84.96875286102295 seconds ---

In [391]: `# Checking the df_result dataframe which contains consolidated results of all the runs`
`df_Results`

Out[391]:

| | Methodology | Model | Accuracy | roc_value | threshold |
|----|---|--|----------|-----------|-----------|
| 0 | RepeatedKFold Cross Validation | Logistic Regression with L2 Regularisation | 0.9990 | 0.950886 | 0.030690 |
| 1 | RepeatedKFold Cross Validation | Logistic Regression with L1 Regularisation | 0.9991 | 0.933580 | 0.057833 |
| 2 | RepeatedKFold Cross Validation | KNN | 0.9990 | 0.866296 | 0.200000 |
| 3 | RepeatedKFold Cross Validation | Tree Model with gini criteria | 0.9989 | 0.833033 | 1.000000 |
| 4 | RepeatedKFold Cross Validation | Tree Model with entropy criteria | 0.9992 | 0.899750 | 1.000000 |
| 5 | RepeatedKFold Cross Validation | Random Forest | 0.9997 | 0.965381 | 0.300000 |
| 6 | RepeatedKFold Cross Validation | XGBoost | 0.9996 | 0.936365 | 0.283122 |
| 7 | RepeatedKFold Cross Validation | SVM | 0.9981 | 0.724487 | 0.003482 |
| 8 | RepeatedKFold Cross Validation | ANN | 0.9984 | 0.927278 | 0.000521 |
| 9 | StratifiedKFold Cross Validation | Logistic Regression with L2 Regularisation | 0.9992 | 0.956638 | 0.002113 |
| 10 | StratifiedKFold Cross Validation | Logistic Regression with L1 Regularisation | 0.9991 | 0.892559 | 0.301463 |
| 11 | StratifiedKFold Cross Validation | KNN | 0.9990 | 0.823250 | 0.200000 |
| 12 | StratifiedKFold Cross Validation | Tree Model with gini criteria | 0.9985 | 0.823079 | 1.000000 |
| 13 | StratifiedKFold Cross Validation | Tree Model with entropy criteria | 0.9990 | 0.823329 | 1.000000 |
| 14 | StratifiedKFold Cross Validation | Random Forest | 0.9995 | 0.937102 | 0.010000 |
| 15 | StratifiedKFold Cross Validation | XGBoost | 0.9993 | 0.971463 | 0.001395 |
| 16 | StratifiedKFold Cross Validation | SVM | 0.9975 | 0.634779 | 0.003424 |
| 17 | StratifiedKFold Cross Validation | ANN | 0.9985 | 0.925528 | 0.000077 |
| 18 | Random Oversampling with StratifiedKFold CV | Logistic Regression with L2 Regularisation | 0.9808 | 0.963078 | 0.383737 |
| 19 | Random Oversampling with StratifiedKFold CV | Logistic Regression with L1 Regularisation | 0.9983 | 0.500000 | 1.500000 |
| 20 | Random Oversampling with StratifiedKFold CV | KNN | 0.9986 | 0.823164 | 0.600000 |
| 21 | Random Oversampling with StratifiedKFold CV | Tree Model with gini criteria | 0.9987 | 0.793817 | 1.000000 |
| 22 | Random Oversampling with StratifiedKFold CV | Tree Model with entropy criteria | 0.9991 | 0.823379 | 1.000000 |
| 23 | Random Oversampling with StratifiedKFold CV | Random Forest | 0.9994 | 0.937891 | 0.020000 |
| 24 | Random Oversampling with StratifiedKFold CV | XGBoost | 0.9992 | 0.971475 | 0.135445 |
| 25 | Random Oversampling with StratifiedKFold CV | ANN | 0.9991 | 0.946933 | 0.000007 |
| 26 | SMOTE Oversampling with StratifiedKFold CV | Logistic Regression with L2 Regularisation | 0.9700 | 0.963647 | 0.370725 |
| 27 | SMOTE Oversampling with StratifiedKFold CV | Logistic Regression with L1 Regularisation | 0.9983 | 0.500000 | 1.500000 |
| 28 | SMOTE Oversampling with StratifiedKFold CV | KNN | 0.9925 | 0.850154 | 1.000000 |
| 29 | SMOTE Oversampling with StratifiedKFold CV | Tree Model with gini criteria | 0.9955 | 0.704129 | 1.000000 |
| 30 | SMOTE Oversampling with StratifiedKFold CV | Tree Model with entropy criteria | 0.9974 | 0.822528 | 1.000000 |
| 31 | SMOTE Oversampling with StratifiedKFold CV | Random Forest | 0.9995 | 0.958948 | 0.180000 |
| 32 | SMOTE Oversampling with StratifiedKFold CV | XGBoost | 0.9948 | 0.946403 | 0.560311 |
| 33 | SMOTE Oversampling with StratifiedKFold CV | ANN | 0.9941 | 0.944962 | 0.060091 |
| 34 | ADASYN Oversampling with StratifiedKFold CV | Logistic Regression with L2 Regularisation | 0.9567 | 0.964887 | 0.535069 |
| 35 | ADASYN Oversampling with StratifiedKFold CV | Logistic Regression with L1 Regularisation | 0.9983 | 0.500000 | 1.500000 |
| 36 | ADASYN Oversampling with StratifiedKFold CV | KNN | 0.9919 | 0.850080 | 1.000000 |
| 37 | ADASYN Oversampling with StratifiedKFold CV | Tree Model with gini criteria | 0.9971 | 0.704931 | 1.000000 |
| 38 | ADASYN Oversampling with StratifiedKFold CV | Tree Model with entropy criteria | 0.9970 | 0.792966 | 1.000000 |
| 39 | ADASYN Oversampling with StratifiedKFold CV | Random Forest | 0.9993 | 0.959888 | 0.150000 |
| 40 | ADASYN Oversampling with StratifiedKFold CV | XGBoost | 0.9941 | 0.952767 | 0.575517 |
| 41 | ADASYN Oversampling with StratifiedKFold CV | ANN | 0.9980 | 0.953533 | 0.001770 |

Results for ADASYN Oversampling with StratifiedKFold:

- Looking at Accuracy and ROC value we have XGBoost which has provided best results for ADASYN Oversampling with StratifiedKFold technique

Overall conclusion after running the models on Oversampled data:

- Looking at above results it seems XGBOOST model with Random Oversampling with StratifiedKFold CV has provided the best results under the category of all oversampling techniques. So we will try to tune the hyperparameters of this model to get best results.

Hyperparameter Tuning

HPT - Xgboost Regression

- Hyperparameter tuning in XGBoost regression involves adjusting the values of certain parameters to optimize the performance of the model on a given dataset.
- Commonly tuned hyperparameters in XGBoost regression include the learning rate, the number of trees in the model, the maximum depth of each tree, the minimum child weight, the subsample ratio, and the regularization coefficients. The optimal values for these hyperparameters can be found using techniques such as grid search, random search, or bayesian optimization.

```
In [392]: # Performing Hyperparameter tuning
from xgboost.sklearn import XGBClassifier
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV

param_test = {
    'max_depth': range(3,10,2),
    'min_child_weight': range(1,6,2),
    'n_estimators': range(60,130,150),
    'learning_rate': [0.05, 0.1, 0.125, 0.15, 0.2],
    'gamma': [1/10.0 for i in range(0,5)],
    'subsample': [1/10.0 for i in range(7,10)],
    'colsample_bytree': [1/10.0 for i in range(7,10)]
}

gsearch1 = RandomizedSearchCV(estimator = XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
    colsample_bynode=1, max_delta_step=0,
    missing=None, n_jobs=-1,
    nthread=None, objective='binary:logistic', random_state=42,
    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
    silent=None, verbosity=1),
    param_distributions = param_test, n_iter=5, scoring='roc_auc', n_jobs=-1, cv=5)

gsearch1.fit(X_over, y_over)
gsearch1.cv_results_, gsearch1.best_params_, gsearch1.best_score_
```

```
Out[392]: ({'mean_fit_time': array([ 7.95628734, 17.85095525,  7.41580715, 13.04083924, 17.3790236 ]),
'std_fit_time': array([0.84154201, 0.32982394, 0.63760971, 0.41301149, 1.98682468]),
'mean_score_time': array([0.05614347, 0.06945863, 0.05744319, 0.06470861, 0.05429459]),
'std_score_time': array([0.01722913, 0.01862621, 0.02510229, 0.01841114, 0.00632305]),
'param_subsample': masked_array(data=[0.9, 0.9, 0.7, 0.7, 0.8],
    mask=[False, False, False, False, False],
    fill_value='?',
    dtype=object),
'param_n_estimators': masked_array(data=[60, 60, 60, 60, 60],
    mask=[False, False, False, False, False],
    fill_value='?',
    dtype=object),
'param_min_child_weight': masked_array(data=[1, 3, 1, 1, 3],
    mask=[False, False, False, False, False],
    fill_value='?',
    dtype=object),
'param_max_depth': masked_array(data=[3, 9, 3, 5, 9],
    mask=[False, False, False, False, False],
    fill_value='?',
    dtype=object),
'param_learning_rate': masked_array(data=[0.15, 0.1, 0.15, 0.1, 0.1],
    mask=[False, False, False, False, False],
    fill_value='?',
    dtype=object),
'param_gamma': masked_array(data=[0.4, 0.4, 0.3, 0.4, 0.3],
    mask=[False, False, False, False, False],
    fill_value='?',
    dtype=object),
'param_colsample_bytree': masked_array(data=[0.7, 0.8, 0.7, 0.9, 0.8],
    mask=[False, False, False, False, False],
    fill_value='?',
    dtype=object),
'params': [{'subsample': 0.9,
'n_estimators': 60,
'min_child_weight': 1,
'max_depth': 3,
'learning_rate': 0.15,
'gamma': 0.4,
'colsample_bytree': 0.7},
{'subsample': 0.9,
'n_estimators': 60,
'min_child_weight': 3,
'max_depth': 9,
'learning_rate': 0.1,
'gamma': 0.4,
'colsample_bytree': 0.8},
{'subsample': 0.7,
'n_estimators': 60,
'min_child_weight': 1,
'max_depth': 3,
'learning_rate': 0.15,
'gamma': 0.3,
'colsample_bytree': 0.7},
{'subsample': 0.7,
'n_estimators': 60,
'min_child_weight': 1,
'max_depth': 5,
'learning_rate': 0.1,
'gamma': 0.4,
'colsample_bytree': 0.9},
{'subsample': 0.8,
'n_estimators': 60,
'min_child_weight': 3,
'max_depth': 9,
'learning_rate': 0.1,
'gamma': 0.3,
'colsample_bytree': 0.8}],
'split0_test_score': array([0.99983833, 0.99996529, 0.99982739, 0.99993807, 0.99997238]),
'split1_test_score': array([0.99991597, 1. , 0.99992885, 0.99997357, 1. ]),
'split2_test_score': array([0.99972677, 0.99999417, 0.99979866, 0.99990744, 0.99997796]),
'split3_test_score': array([0.99989202, 0.99992623, 0.99987478, 0.99988858, 0.99992124]),
'split4_test_score': array([0.99993421, 1. , 0.99992461, 0.99980818, 0.99998206]),
'mean_test_score': array([0.99986146, 0.99997714, 0.99987086, 0.99993757, 0.99997073]),
'std_test_score': array([7.46656402e-05, 2.85217757e-05, 5.17106231e-05, 3.58332424e-05,
    2.64156355e-05]),
'rank_test_score': array([5, 1, 4, 3, 2], dtype=int32)},
{'subsample': 0.9,
'n_estimators': 60,
'min_child_weight': 3,
'max_depth': 9,
'learning_rate': 0.1,
'gamma': 0.4,
'colsample_bytree': 0.8},
0.9999771372442581)
```

- Please note that the hyperparameters found above using RandomizedSearchCV and the hyperparameters used below in creating the final model might be different, the reason being, I have executed the RandomizedSearchCV multiple times to find which set of hyperparameters gives the optimum result and finally used the one below which gave me the best performance.


```
In [393]: # Creating XGBoost model with selected hyperparameters
from xgboost import XGBClassifier

clf = XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                    colsample_bynode=1, colsample_bytree=0.7, gamma=0.2,
                    learning_rate=0.125, max_delta_step=0, max_depth=7,
                    min_child_weight=5, missing=None, n_estimators=60, n_jobs=1,
                    nthread=None, objective='binary:logistic', random_state=42,
                    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                    silent=None, subsample=0.8, verbosity=1)

# fit on the dataset
clf.fit(X_over, y_over)
XGB_test_score = clf.score(X_test, y_test)
print('Model Accuracy: {}'.format(XGB_test_score))

# Probabilities for each class
XGB_probs = clf.predict_proba(X_test)[: , 1]

# Calculate roc auc
XGB_roc_value = roc_auc_score(y_test, XGB_probs)

print("XGboost roc_value: {}".format(XGB_roc_value))
fpr, tpr, thresholds = metrics.roc_curve(y_test, XGB_probs)
threshold = thresholds[np.argmax(tpr-fpr)]
print("XGBoost threshold: {}".format(threshold))

Model Accuracy: 0.9991
XGboost roc_value: 0.9680280005420979
XGBoost threshold: 0.0020904664415866137
```

Print the important features of the best model to understand the dataset

```
In [394]: imp_var = []
for i in clf.feature_importances_:
    imp_var.append(i)
print('Top var =', imp_var.index(np.sort(clf.feature_importances_)[-1])+1)
print('2nd Top var =', imp_var.index(np.sort(clf.feature_importances_)[-2])+1)
print('3rd Top var =', imp_var.index(np.sort(clf.feature_importances_)[-3])+1)


Top var = 14
2nd Top var = 17
3rd Top var = 10
```

```
In [395]: # Calculate roc auc
XGB_roc_value = roc_auc_score(y_test, XGB_probs)

print("XGboost roc_value: {}".format(XGB_roc_value))
fpr, tpr, thresholds = metrics.roc_curve(y_test, XGB_probs)
threshold = thresholds[np.argmax(tpr-fpr)]
print("XGBoost threshold: {}".format(threshold))

XGboost roc_value: 0.9680280005420979
XGBoost threshold: 0.0020904664415866137
```

```
In [396]: # Confusion matrix
Plot_confusion_matrix(y_test,clf.predict(X_test))
```



Conclusion

- In the oversample cases, of all the models we build found that the XGBOOST model with Random Oversampling with StratifiedKFold CV gave us the best accuracy and ROC on oversampled data. Post that we performed hyperparameter tuning and got the below metrics :
 - XGboost roc_value: 0.9680280005420979
 - XGBoost threshold: 0.0020904664415866137
- However, of all the models we created we found Logistic Regression with L2 Regularisation for StratifiedKFold cross validation (without any oversampling or undersampling) gave us the best result.
- Also, all the models we created gave good results